


Advanced Database- IS411




Introduced by

Dr. Ebtsam Adel

Lecturer of Information Systems,
Information Systems department,
Faculty of computers and information,
Damanhour university



Materials





1. Which of the following properties does a transaction not have?

- a. Atomicity
- b. Consistency
- c. Isolation
- d. Scalability

2. What does the 'Durability' property ensure in a transaction?

- a. Transactions are completed without errors
- b. Once a transaction is committed, it remains so even in case of a failure
- c. Transactions can be rolled back
- d. Transactions are isolated from each other

3. Which property of transactions guarantees that concurrent transactions do not interfere with each other?

- a. Atomicity
- b. Consistency
- c. Isolation
- d. Durability

4. What is the primary purpose of the commit operation in a transaction?

- a. To start a transaction
- b. To save changes made during the transaction
- c. To undo changes made
- d. To check for errors in the transaction



Topics

- ✓ **Chapter 20** Introduction to Transaction Processing Concepts and Theory
- ✓ **chapter 24** NOSQL Databases and Big Data Storage Systems
- ✓ **chapter 25** Big Data Technologies Based on MapReduce and Hadoop
- ✓ **chapter 27** Introduction to Information Retrieval and Web Search
- ✓ **chapter 29** Overview of Data Warehousing and OLAP
- ✓ **chapter 30** Database Security

CHAPTER 24

NOSQL Databases and Big Data Storage Systems



Data types

- ❑ Data could be presented in three different types of data: **structured**, **semi-structured**, and **unstructured** data.
- ❑ **Structured** data as (**database**), **semi-structured** data (**XML files**, **JSON documents**), and **unstructured** data such as (**documents**, **video**, e-mails, and **reports**).

Unstructured data

The university has 5600 students.
John's ID is number 1, he is 18 years old and already holds a B.Sc. degree.
David's ID is number 2, he is 31 years old and holds a Ph.D. degree. Robert's ID is number 3, he is 51 years old and also holds the same degree as David, a Ph.D. degree.

Semi-structured data

```
<University>
  <Student ID="1">
    <Name>John</Name>
    <Age>18</Age>
    <Degree>B.Sc.</Degree>
  </Student>
  <Student ID="2">
    <Name>David</Name>
    <Age>31</Age>
    <Degree>Ph.D. </Degree>
  </Student>
  ....
</University>
```

Structured data

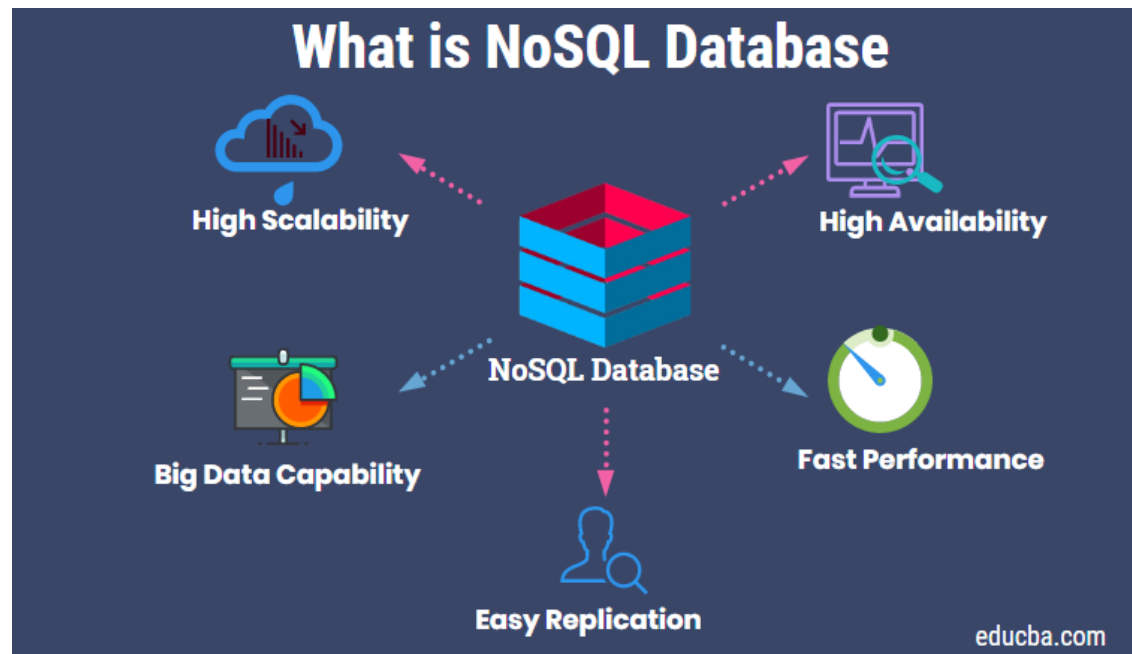
ID	Name	Age	Degree
1	John	18	B.Sc.
2	David	31	Ph.D.
3	Robert	51	Ph.D.
4	Rick	26	M.Sc.
5	Michael	19	B.Sc.

Introduction

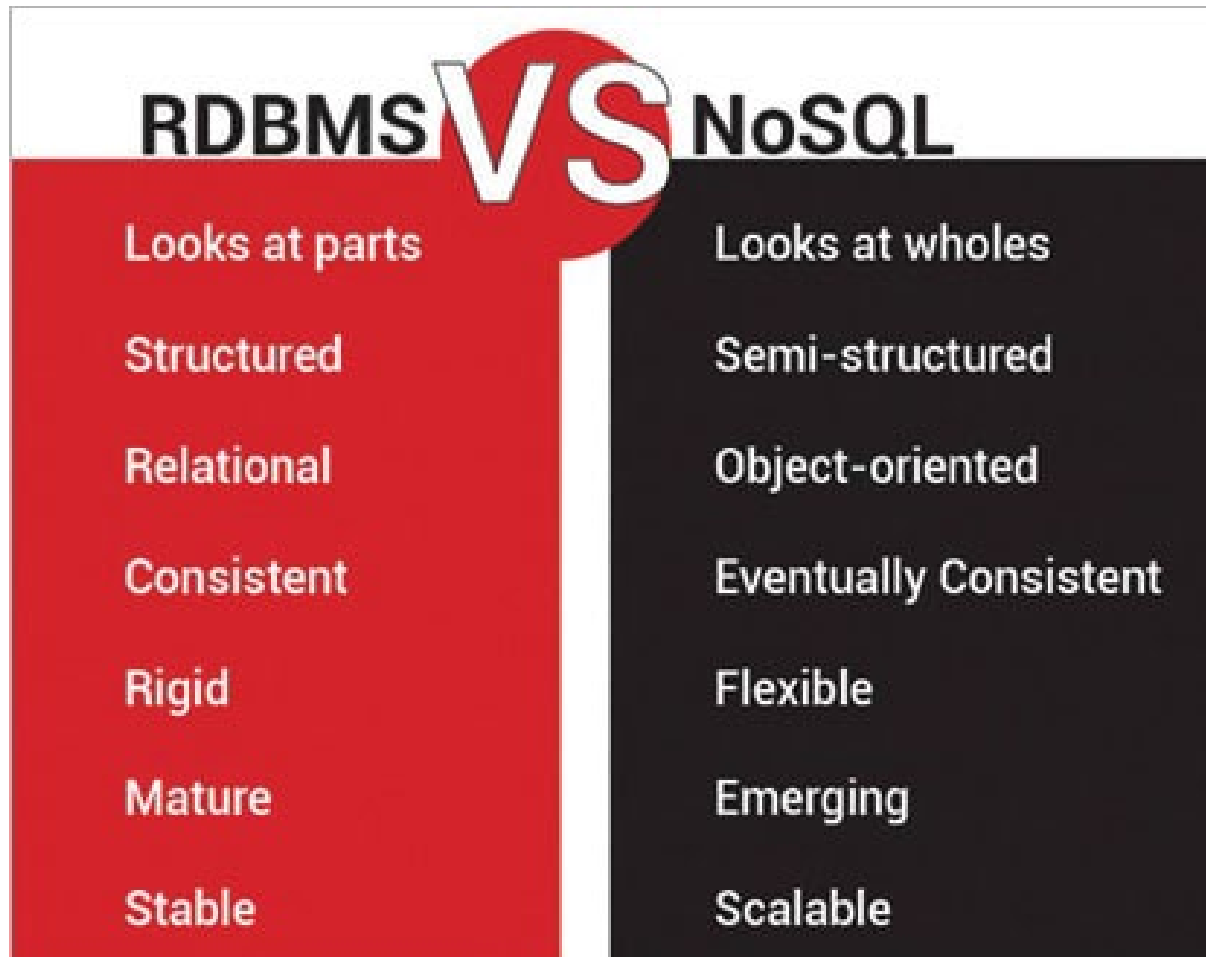
- **Most SQL databases** are relational. Relational Databases are **tabular** and have a **pre-defined schema** that organizes the data logically.
- ❑ **NoSQL** is a **non-relational** database management system for certain data models. These data models **don't need a schema** and are scalable.
- ❑ **NoSQL databases** don't require a certain schema. They prioritize speed and flexibility in data storage. **Amazon, Facebook, and Google.**

Introduction

- NOSQL: **Not only SQL**
- Most NOSQL systems are **distributed databases** or distributed storage systems.
 - Focus on **semi-structured data storage**, high performance, **availability**, **data replication**, and **scalability**.



Introduction



Introduction (cont'd.)


- NOSQL systems focus on storage of “big data”.
- Typical applications that use NOSQL
 - Social media
 - Web links
 - User profiles
 - Marketing and sales
 - Posts and tweets
 - Road maps and spatial data
 - Email
 -

24.1 Introduction to NOSQL Systems

- BigTable
 - Google's proprietary NOSQL system
 - Column-based or wide column store
- DynamoDB (Amazon)
 - Key-value data store
- Cassandra (Facebook)
 - Uses concepts **from both** key-value store and column-based systems.



Wide-Column Store



Introduction to NOSQL Systems (cont'd.)

- MongoDB and CouchDB



- Document stores

- Neo4J and GraphBase

- Graph-based NOSQL systems

- OrientDB



- Combines several concepts

- Database systems classified on the object model


- Or native XML model

NOSQL characteristics related to *distributed databases*

- NOSQL characteristics related to **distributed databases** and distributed systems
 - Scalability
 - Availability, replication, and eventual consistency
 - Replication models
 - Master-slave
 - Master-master
 - Sharding of files
 - High performance data access

Master-Slave Replication


- The process of **copying and synchronizing** data from a primary database (the **master**) to one or more secondary databases (the **slaves**) is known as master-slave replication.
- In this configuration, all write operations, including inserts, updates, and deletions, **must be received** by the master database.
- After modifications are made to the master database, a copy of the data is kept in the slave databases.



Master-Slave Replication

Master-Master Replication

- is a configuration where two or more databases are set up as **master databases**, each of which is able to **accept write** operations.
- In other words, any modifications made to one master database are **reflected** in all other master databases within the setup.



Master-Master Replication

NOSQL characteristics related to *data models* and *query languages*

NOSQL characteristics related to **data models** and **query languages**

- ✓ Less powerful query languages
- ✓ Versioning
- ✓ Schema not required

Enter MySQL Query:

```
1 SELECT Type FROM Places
2 WHERE Type IN('Type1', 'Type 2')
3 ORDER BY Type;
```

MongoDB Syntax:

```
1 db.Places.find({
2     "Type": {
3         "$in": ["Type1", "Type 2"]
4     }
5 }, {
6     "Type": 1
7 }).sort({
8     "Type": 1
9 });
```

Introduction to NOSQL Systems (cont'd.)

- **Versioning**: multiple versions of the data items, with the timestamps of when the data version was created.

Introduction to NOSQL Systems (cont'd.)

Categories of NOSQL systems

- Document-based NOSQL systems
- NOSQL key-value stores
- Column-based or wide column NOSQL systems
- Graph-based NOSQL systems
- Hybrid NOSQL systems
- Object databases
- XML databases

The CAP Theorem

24.2 The CAP Theorem

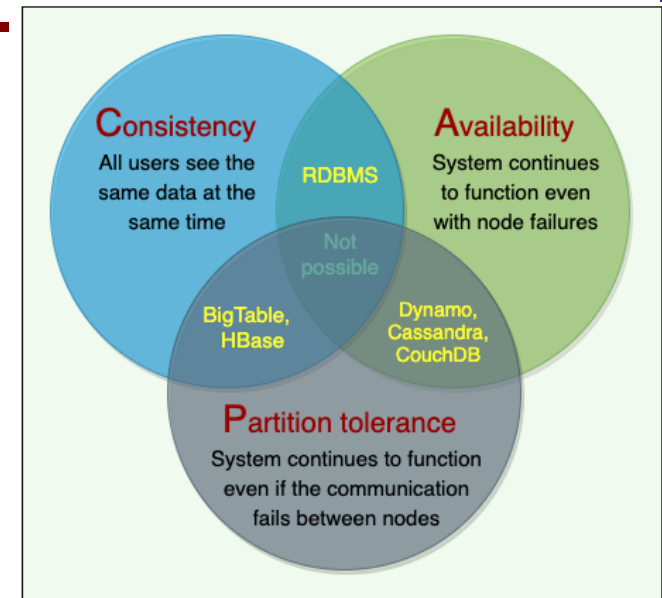
- ❑ **Consistency** means that the nodes will have the same copies of a replicated data item visible for various transactions.
- ❑ **Availability** means that each read or write request for a data item will either be processed successfully or will receive a message that the operation cannot be completed.
- ❑ **partition tolerance** (in the face of the nodes in the system being partitioned by a network fault).

24.2 The CAP Theorem

- Various levels of consistency among replicated data items
 - Enforcing **serializability** the strongest form of consistency
 - High **overhead** – can reduce read/write operation performance
- CAP theorem
 - Consistency, availability, and partition tolerance
 - **Not possible to guarantee** all three simultaneously
 - In distributed system with data replication


The CAP Theorem (cont'd.)

- Designer can choose two of three to guarantee
 - Weaker consistency level is often acceptable in NOSQL distributed data store
 - Guaranteeing availability and partition tolerance more important
 - Eventual consistency often adopted.



Types of NoSQL Databases


Key Value




Graph DB



Column Family



Document




Types of NoSQL Databases

SQL Database

Relational




Analytical (OLAP)




NoSQL Database


Column-Family




Graph




Document



Key-Value



Document-Based NOSQL



Document
Store

24.3 Document-Based NOSQL Systems and MongoDB

- Document stores
 - Collections of similar documents
- Individual documents resemble complex objects or XML documents
 - Documents are self-describing
 - Can have different data elements
- Documents can be specified in various formats
 - XML
 - **JSON**
 - **BSON**

MongoDB Data Model

- Documents stored in binary JSON (BSON) format
- Individual documents stored in a collection
- Example command
 - First parameter specifies name of the collection
 - Collection options include limits on size and number of documents

```
db.createCollection("project", { capped : true, size : 1310720, max : 500 } )
```

- Each document in collection has unique ObjectID field called **_id**

MongoDB Data Model (cont'd.)

- A collection **does not have a schema**
 - Structure of the data fields in documents chosen based on how documents will be accessed
 - User can choose normalized or denormalized design
- Document creation using **insert operation**

```
db.<collection_name>.insert(<document(s)>)
```

- Document deletion using **remove operation**

```
db.<collection_name>.remove(<condition>)
```


MongoDB Distributed Systems Characteristics

- **Two-phase commit** method
 - Used to ensure atomicity and consistency of multidocument transactions
- Replication in MongoDB
 - Concept of replica set to **create multiple copies** on different nodes
 - Variation of **master-slave approach**
 - Primary copy, secondary copy, and arbiter
 - Arbiter participates in **elections** to **select new primary** if needed


MongoDB Distributed Systems Characteristics (cont'd.)

- **Replication** in MongoDB (cont'd.)
 - All write operations applied to the primary copy and propagated to the secondaries
 - User can choose **read preference**
 - Read requests can be processed at any replica
- **Sharding** in MongoDB
 - Horizontal partitioning divides the documents into disjoint partitions (**shards**)
 - Allows adding more nodes as needed
 - Shards stored on different nodes to achieve **load balancing**

MongoDB Distributed Systems Characteristics (cont'd.)

- Sharding in MongoDB (cont'd.)
 - Partitioning field (**shard key**) must exist in every document in the collection
 - Must have an **index**
 - Range partitioning
 - Creates **chunks** by specifying a range of key values
 - Works best with **range queries**
 - Hash partitioning
 - Partitioning based on the **hash values** of each shard key

NOSQL Key-Value



Key-Value
Store

24.4 NOSQL Key-Value Stores

- Key-value stores focus on high performance, availability, and scalability.
 - Can store structured, unstructured, or semi-structured data
- **Key**: unique identifier associated with a data item
 - Used for fast retrieval
- **Value**: the data item itself
 - Can be string or array of bytes
 - Application interprets the structure
- Not support complex queries.

DynamoDB Overview

- **DynamoDB** part of Amazon's Web Services/SDK platforms
- Table holds a collection of **self-describing items**
- Item consists of **attribute-value pairs**
 - Attribute values can be single or multi-valued
- Primary key used to locate items within a table
 - Can be single attribute or pair of attributes

Voldemort Key-Value Distributed Data Store

- Voldemort: open source key-value system similar to **DynamoDB**
- Voldemort features
 - Simple basic operations (get, put, and delete)
 - High-level formatted data values
 - Consistent **hashing** for distributing (key, value) pairs
 - Consistency and versioning
 - Concurrent writes allowed
 - Each write associated with a vector clock

Examples of Other Key-Value Stores

- Oracle key-value store
 - Oracle NOSQL Database

Ex. Redis

- Caches data in main memory to improve performance
 - Offers master-slave replication and high availability
 - Offers persistence by backing up cache to disk.

Column-Based NoSQL databases



Wide-Column
Store

24.5 Column-Based or Wide Column NOSQL Systems

- Data is stored in **columns** rather than **rows**, enabling high-speed **analytics** and **distributed** computing.
- Efficient for **handling large-scale data** with high write/read demands.
- Great for **time-series data**, **IoT applications**, and big data analytics.
- **Examples:** **Apache Cassandra**, **HBase**, **Google Bigtable**.

24.5 Column-Based or Wide Column NOSQL Systems

row-store



+ easy to add/modify a record

- might read in unnecessary data

column-store



+ only need to read in relevant data

- tuple writes require multiple accesses

=> suitable for read-mostly, read-intensive, large data repositories

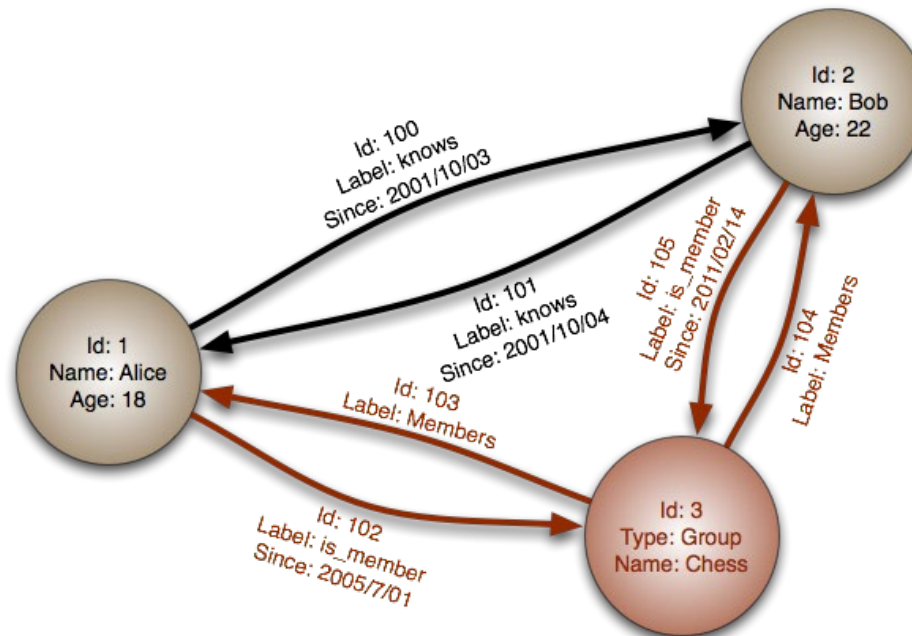
Hbase Data Model and Versioning (cont'd.)

- HBase stores multiple versions of data items
 - Timestamp associated with each version
- Each row in a table has a **unique row key**
- Table associated with **one or more column families**
- Column qualifiers can be dynamically specified as new table rows are created and inserted
- **Namespace** is **collection of tables**
- Cell holds a basic data item

Hbase Crud Operations

- Provides only low-level **CRUD** (create, read, update, delete) operations.
- Application programs implement more complex operations
- Create
 - Creates a new table and specifies one or more column families associated with the table
- Put
 - Inserts new data or new versions of existing data items

NOSQL Graph



24.6 NOSQL Graph Databases and Neo4j

■ Graph databases

- Data represented as a **graph**
- Collection of vertices (**nodes**) and **edges**
- Possible to store data associated with both individual nodes and individual edges.
- Can perform relations faster than SQL

■ Example: Neo4j, VertexDB

Neo4j (cont'd.)

- Nodes can have labels
 - Zero, one, or several
- Both nodes and relationships can have properties
- Each relationship has a start node, end node, and a relationship type
- Properties specified using a map pattern
- Somewhat similar to ER/EER concepts

Neo4j (cont'd.)

- Creating nodes in Neo4j
 - CREATE command
 - Part of high-level declarative query language
 - Node label can be specified when node is created
 - Properties are enclosed in curly brackets

24.7 Summary

- NOSQL systems focus on storage of “big data”
- General categories
 - Document-based
 - Key-value stores
 - Column-based
 - Graph-based
 - Some systems use techniques spanning two or more categories
- Consistency paradigms
- CAP theorem

Thank You!

Any questions? 