


# Advanced Database- IS411



Introduced by

**Dr. Ebtsam Adel**

Lecturer of Information Systems,  
Information Systems department,  
Faculty of computers and information,  
Damanhour university



# Lecture (02)

---

## **Chapter 20: Introduction to Transaction Processing Concepts and Theory**

Part 2

Slide 1-2

# Outline

---

- ❑ Characterizing Schedules based on Recoverability
- ❑ Characterizing Schedules based on Serializability

---

# **Characterizing Schedules Based on Recoverability**

# Schedules of Transactions

- **Transaction schedule (history)**: When transactions are executing **concurrently** in an **interleaved** fashion, then the **order of execution** of operations from the various transactions forms what is known as a **transaction schedule (history)**.
- The following figure shows 4 possible schedules (A, B, C, D) of two transactions T1 and T2:
  - ✓ Order of operations from top to bottom
  - ✓ Each schedule includes same operations
  - ✓ Different order of operations in each schedule

# Schedules of Transactions (cont.)

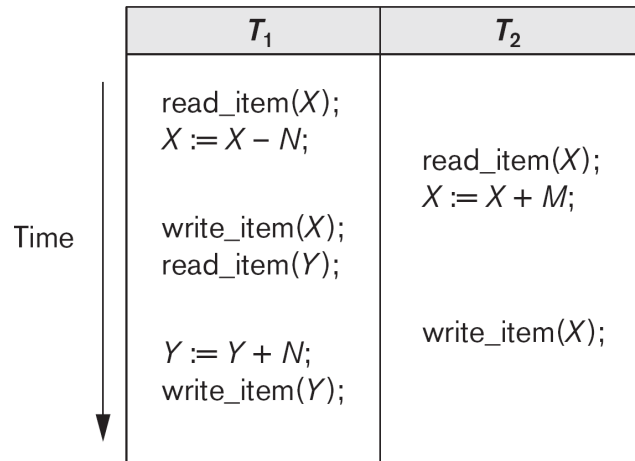
- **Formal definition** of a **schedule** (or **history**) **S** of  $n$  transactions  $T_1, T_2, \dots, T_n$  is an **ordering** of all the operations of the transactions subject to the constraint that, for each transaction  $T_i$  that participates in  $S$ , the operations of  $T_i$  in  $S$  must **appear** in **the same order** in which they occur in  $T_i$ .
- The order of operations in **S** is considered to be a **total ordering**, meaning that for any two operations in the schedule, **one must occur before the other**.

# Schedules of Transactions (cont.)

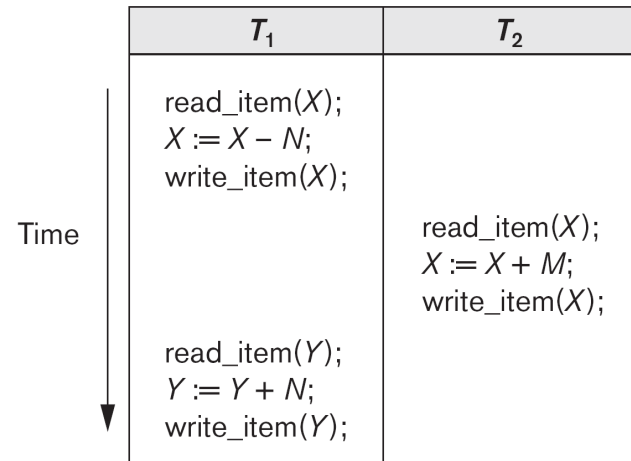
---

- Schedules can also be displayed in more compact notation.
- Order of operations from **left** to **right**.
- Include only **read** (r) and **write** (w) operations, with **transaction id** (1, 2, ...) and **item name** (X, Y, ...)
- Can also include other operations such as b (**begin**), e (**end**), c (**commit**), a (**abort**).

# Schedules of Transactions (cont.)



Schedule C



Schedule D

**Schedule C:** r1(X); r2(X); w1(X); r1(Y); w2(X); w1(Y);

**Schedule D:** r1(X); w1(X); r2(X); w2(X); r1(Y); w1(Y);



# Schedules of Transactions (cont.)

$T_1$	$T_2$
read_item(X); $X := X - N$ ; write_item(X); read_item(Y); $Y := Y + N$ ; write_item(Y);	read_item(X); $X := X + M$ ; write_item(X);

Schedule A

$T_1$	$T_2$
read_item(X); $X := X - N$ ; write_item(X); read_item(Y); $Y := Y + N$ ; write_item(Y);	read_item(X); $X := X + M$ ; write_item(X);

Schedule B

**Schedule A:** r1(X); w1(X); r1(Y); w1(Y); r2(X); w2(x);

**Schedule B:** r2(X); w2(X); r1(X); w1(X); r1(Y); w1(Y);

# Schedules of Transactions- Conflict

**Conflicting Operations in a Schedule.** Two operations in a schedule are said to **conflict** if they satisfy all three of the following conditions: (1) they belong to *different transactions*; (2) they access the *same item X*; and (3) *at least one* of the operations is a write\_item(X). For example, in schedule  $S_a$ , the operations  $r_1(X)$  and  $w_2(X)$  conflict, as do the operations  $r_2(X)$  and  $w_1(X)$ , and the operations  $w_1(X)$  and  $w_2(X)$ . However, the operations  $r_1(X)$  and  $r_2(X)$  do not conflict, since they are both read

$S_a: r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); w_1(Y);$

Conflict operations	
R1(x)	W2(x)
R2(x)	W1(x)
W1(x)	W2(x)

Check the conflict.

# Schedules of Transactions- **complete schedule**

---

- A schedule  $S$  of  $n$  transactions  $T_1, T_2, \dots, T_n$  is said to be a **complete schedule** if the following conditions hold:
  1. The operations in  $S$  are exactly those operations in  $T_1, T_2, \dots, T_n$ , including a **commit** or **abort** operation as the **last operation** for each transaction in the schedule.
  2. For any pair of operations from the same transaction  $T_i$ , their relative **order** of appearance in  $S$  is the **same** as their **order** of appearance in  $T_i$ .
  3. For any two **conflicting** operations, one of the two must occur **before** the other in the schedule.

---

# Characterizing Schedules based on Recoverability

# Characterizing Schedules based on Recoverability

---

- First, we would like to ensure that, once a transaction  $T$  is **committed**, it should **never** be necessary to **roll back**  $T$ . This ensures that the **durability property** of transactions is not violated.
- The **schedules** that theoretically meet this criteria are called **recoverable schedules**; Those that do not are called **nonrecoverable** and hence should not be **permitted** by the DBMS.

# Characterizing Schedules based on Recoverability

---

## Schedules classified into two main classes:

- **Recoverable schedule:** One where no committed transaction needs to be rolled back (aborted).  
A schedule  $S$  is **recoverable** if no transaction  $T$  in  $S$  commits **until** all transactions  $T'$  that have **written** an item that  $T$  reads have **committed**.
- **Non-recoverable schedule:** A schedule where a committed transaction may have to be rolled back during recovery.

*Recoverable schedule: Write first must commit first.*

# Characterizing Schedules based on Recoverability

$S_a': r_1(X); r_2(X); w_1(X); r_1(Y); w_2(X); c_2; w_1(Y); c_1;$



**nonrecoverable**

$S_c: r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); c_2; a_1;$

**nonrecoverable**

$S_d: r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); c_1; c_2;$

**recoverable**



$S_e: r_1(X); w_1(X); r_2(X); r_1(Y); w_2(X); w_1(Y); a_1; a_2;$

**nonrecoverable**

# Characterizing Schedules Based on Recoverability (cont.)

---

- Schedule A below is **non-recoverable** because T2 reads the value of X that was written by T1, but then T2 commits before T1 commits or aborts
- To make it **recoverable**, the commit of T2 (c2) must be delayed until T1 either commits, or aborts (Schedule B).
- If T1 commits, T2 can commit
- If T1 aborts, T2 must also abort because it read a value that was written by T1; this value must be undone (reset to its old value) when T1 is aborted
  - known as **cascading rollback**



# Characterizing Schedules based on Recoverability (cont.)

---

- ❑ it is possible for a phenomenon known as **cascading rollback** (or **cascading abort**) to occur in some recoverable schedules, where an **uncommitted** transaction has to be rolled back because it read an item from a transaction that **failed**.
- ❑ A **schedule** is said to be **cascadeless**, or to **avoid cascading rollback**, if every transaction in the schedule **reads only** items that were written by **committed** transactions.

# Characterizing Schedules based on Recoverability (cont.)

---

- ❑ A schedule is said to be **cascadeless**, or to **avoid cascading rollback**, if every transaction in the schedule reads only items that were written by committed transactions. In this case, all items read will not be discarded because the transactions that wrote them have committed, so no cascading rollback will occur.

# Characterizing Schedules based on Recoverability (cont.)

- **Strict schedule:** A schedule in which a transaction T2 can **neither read nor write** an item X until the transaction T1 that last wrote X **has committed**.
- It is important to note that **any strict schedule** is also **cascadeless**, and any **cascadeless** schedule is also **recoverable**.

$S_f: w_1(X, 5); w_2(X, 8); a_1;$

**Cascadeless but  
Unrestricted**

## Example

X= 9

W1(X)= 5    BFIM=9

W2(X)= 8    BFIM=9

Abort T1;

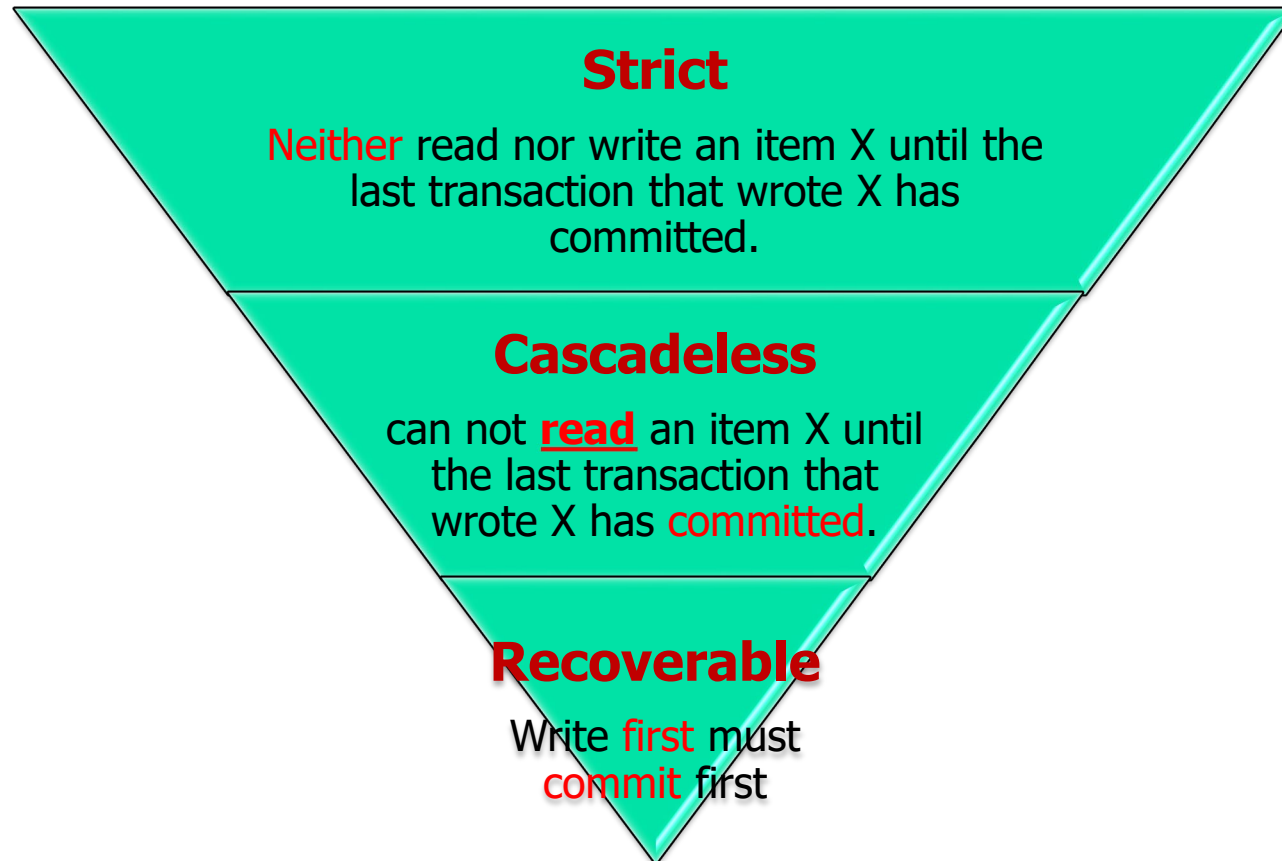
X= 9

**BFIM=Before image**

---

# Exercises

# Exercises



- **Strict "means"** strict, cascadeless, recoverable
- **Cascadeless "means"** cascadeless, recoverable

# Exercises

$S_a$ :  $r1(x); r2(x); \mathbf{w2(x)}; w1(x); \mathbf{c2}; \mathbf{r1(x)}; r1(y); c1 \dots$  **cascadeless**

---

$S_b$ :  $r1(x); r2(x); \mathbf{w1(x)}; r1(y); w1(y); \mathbf{c1, w2(x)}; c2 \dots$  **strict**

---

$S_c$ :  $r1(x); \mathbf{w1(x)}; r1(y); w1(y); \mathbf{r2(x)}; \mathbf{c1}; w2(x); \mathbf{c2} \dots$  **recoverable**


---

$S_d$ :  $r1(x); \mathbf{w1(x)}; r1(y); w1(y); \mathbf{r2(x)}; w2(x); \mathbf{c2}; \mathbf{c1} \dots$

**non\_recoverable**

# Types of schedules in DBMS

---



# Outline

---

- ❑ Characterizing Schedules based on Recoverability
- ❑ Characterizing Schedules based on Serializability



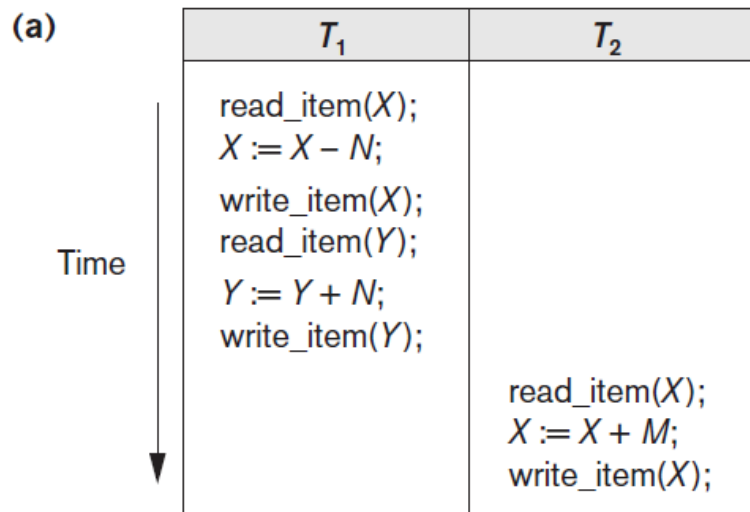
---

# Characterizing Schedules Based on Serializability

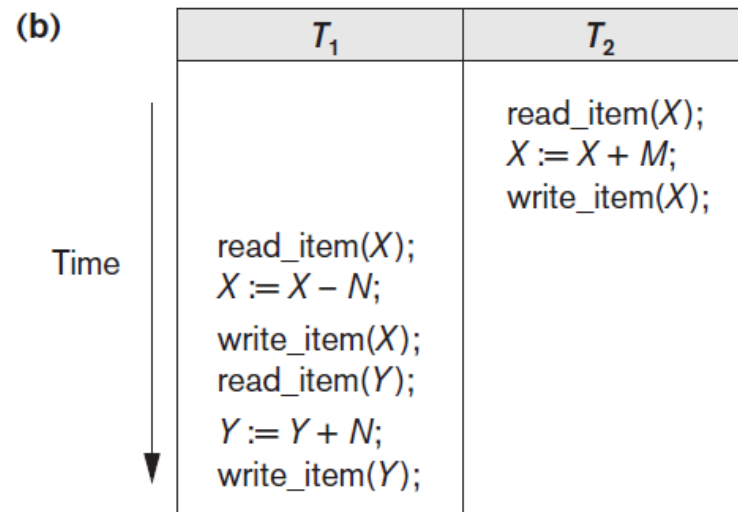
# Characterizing Schedules based on Serializability (cont.)

## 20.5.1 Serial, Nonserial, and Conflict-Serializable Schedules

Schedules A and B in Figures 20.5(a) and (b) are called *serial* because the operations of each transaction are executed consecutively, without any interleaved operations from the other transaction. In a serial schedule, entire transactions are performed in serial order:  $T_1$  and then  $T_2$  in Figure 20.5(a), and  $T_2$  and then  $T_1$  in Figure 20.5(b). Schedules C and D in Figure 20.5(c) are called *nonserial* because each sequence interleaves operations from the two transactions.



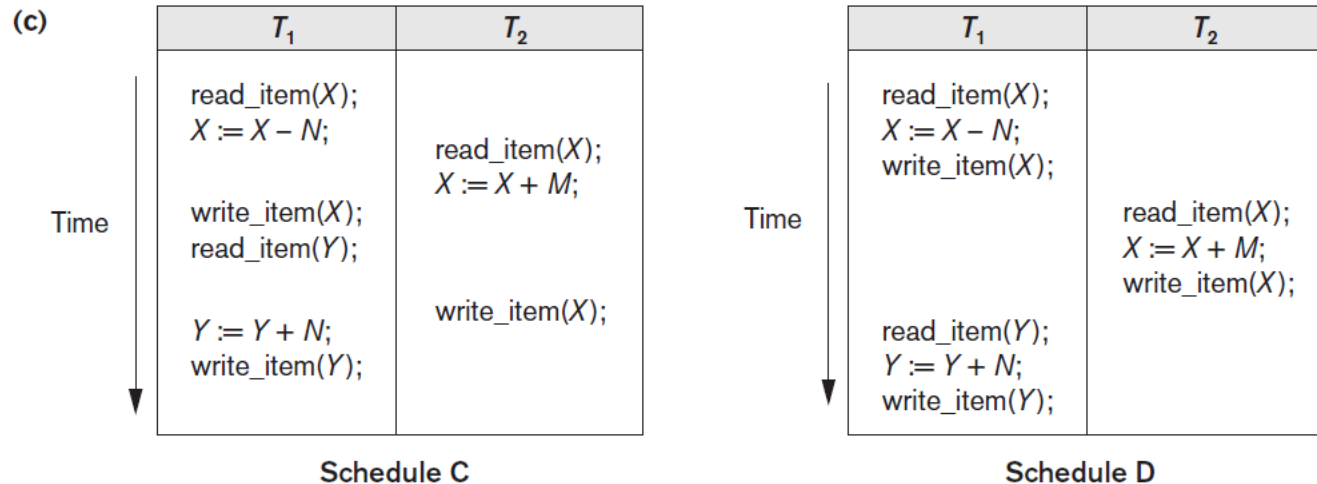
Schedule A



Schedule B

Two serial schedules

# Characterizing Schedules based on Serializability (cont.)



Two non-serial schedules C and D with interleaving of operations.

# Characterizing Schedules based on Serializability (cont.)

---

- **Serial schedule:** A schedule S is **serial** if, for every transaction T participating in the schedule, all the operations of T are executed consecutively (**without interleaving** of operations from other transactions) in the schedule. **Otherwise**, the schedule is called **nonserial**.
- Based on the **consistency preservation** property, any **serial schedule** will produce a **correct result** (assuming no inter-dependencies among different transactions)

# Characterizing Schedules based on Serializability

- ❑ The **problem** with **serial** schedules is that they **limit concurrency** by prohibiting interleaving of operations. In a serial schedule, if a transaction waits for an **I/O operation** to complete, we cannot **switch the CPU processor** to another transaction, thus **wasting valuable CPU processing time**. Additionally, if some transaction T is long, the other transactions must wait for T to complete all its operations before starting. **Hence, serial schedules are unacceptable in practice.**



- ❑ However, if we can determine which other schedules are **equivalent** to a serial schedule, we can allow these schedules to occur.

# Characterizing Schedules based on Serializability (cont.)

---

- Serial schedules are not feasible for performance reasons:
  - No interleaving of operations
  - Long transactions force other transactions to wait
  - System cannot switch to other transaction when a transaction is waiting for disk I/O or any other event.
  - Need to allow concurrency with interleaving without sacrificing correctness

# Characterizing Schedules based on Serializability (cont.)

- **Serializable schedule** (قابل للتسلسل): A schedule  $S$  is **serializable** if it is **equivalent** to some **serial** schedule of the same  $n$  transactions.
- There are  $(n)!$  serial schedules for  $n$  transactions – a **serializable schedule** can be **equivalent** to some serial schedule of the same transactions.
- **Question:** How do we define equivalence of schedules?
  - a. Result equivalent
  - b. Conflict equivalent



# Characterizing Schedules based on Serializability (cont.)

- **Serializable schedule:** A schedule S is **serializable** if it is **equivalent** to some serial schedule of the same n transactions.

**Question:** How do we define **equivalence** of schedules?

- a. Result equivalent
- b. Conflict equivalent





# Equivalence of Schedules- **Result equivalent**

---

- There are several **ways** to define schedule equivalence. The **simplest** but least satisfactory definition involves **comparing the effects** of the schedules on the database.
- ❖ **Result equivalent:** Two schedules are called result equivalent if they produce the **same final state** of the database.

# Equivalence of Schedules- Result equivalent

**Figure 21.6**

Two schedules that are result equivalent for the initial value of  $X = 100$  but are not result equivalent in general.

$S_1$
read_item(X); $X := X + 10$ ; write_item(X);

$S_2$
read_item(X); $X := X * 1.1$ ; write_item (X);

Two schedules are the same result when  $x=100$  **only**.

# Equivalence of Schedules- **Result equivalent**

---

- Difficult to determine without analyzing the **internal operations of the transactions**, which is **not feasible** in general.
- May also get result equivalence by chance for a particular **input parameter** even though schedules are not equivalent in general.



# Equivalence of Schedules - **Conflict equivalent**

---

- **Conflict equivalent:** Two schedules are conflict equivalent if the **order** of any two **conflicting** operations **is the same** in both schedules.
- Commonly used definition of schedule equivalence
- Two operations are **conflicting** if:
  - They access the **same data item X**
  - They are from two **different transactions**
  - At least one is a **write operation**
- Read-Write conflict example:  $r_1(X)$  and  $w_2(X)$
- Write-write conflict example:  $w_1(Y)$  and  $w_2(Y)$

# Equivalence of Schedules - **Conflict equivalent**

S1		S2	
T1	T2	T1	T2
----		W(X)	
	-----		R(X)
W(X)		---	
	R(X)		----

**Conflict equivalent**

S1		S2	
T1	T2	T1	T2
----		----	
	-----		----
W(X)		R(X)	
	R(X)		W(X)

**Not Conflict equivalent**

# Equivalence of Schedules (cont.)

---

- Changing the order of conflicting operations generally causes a **different outcome**.
- **Example:** changing  $r1(X); w2(X)$  to  $w2(X); r1(X)$  means that T1 will read a different value for X
- **Example:** changing  $w1(Y); w2(Y)$  to  $w2(Y); w1(Y)$  means that the final value for Y in the database can be different
- Note that read operations **are not conflicting**; changing  $r1(Z); r2(Z)$  to  $r2(Z); r1(Z)$  does not change the outcome.

# Conflict serializable schedule

---

- ❑ Using the notion of conflict equivalence, we define a schedule  $S$  to be **serializable** if it is (**conflict**) equivalent to some serial schedule  $S'$ .
- ❑ **In such a case**, we can **reorder** the **non-conflicting** operations in  $S$  until we form the **equivalent serial** schedule  $S'$ .

# Conflict serializable schedule

$T_1$	$T_2$
read_item(X); $X := X - N$ ; write_item(X);  read_item(Y); $Y := Y + N$ ; write_item(Y);	read_item(X); $X := X + M$ ; write_item(X);

Schedule D

$T_1$	$T_2$
read_item(X); $X := X - N$ ; write_item(X); read_item(Y); $Y := Y + N$ ; write_item(Y);	read_item(X); $X := X + M$ ; write_item(X);

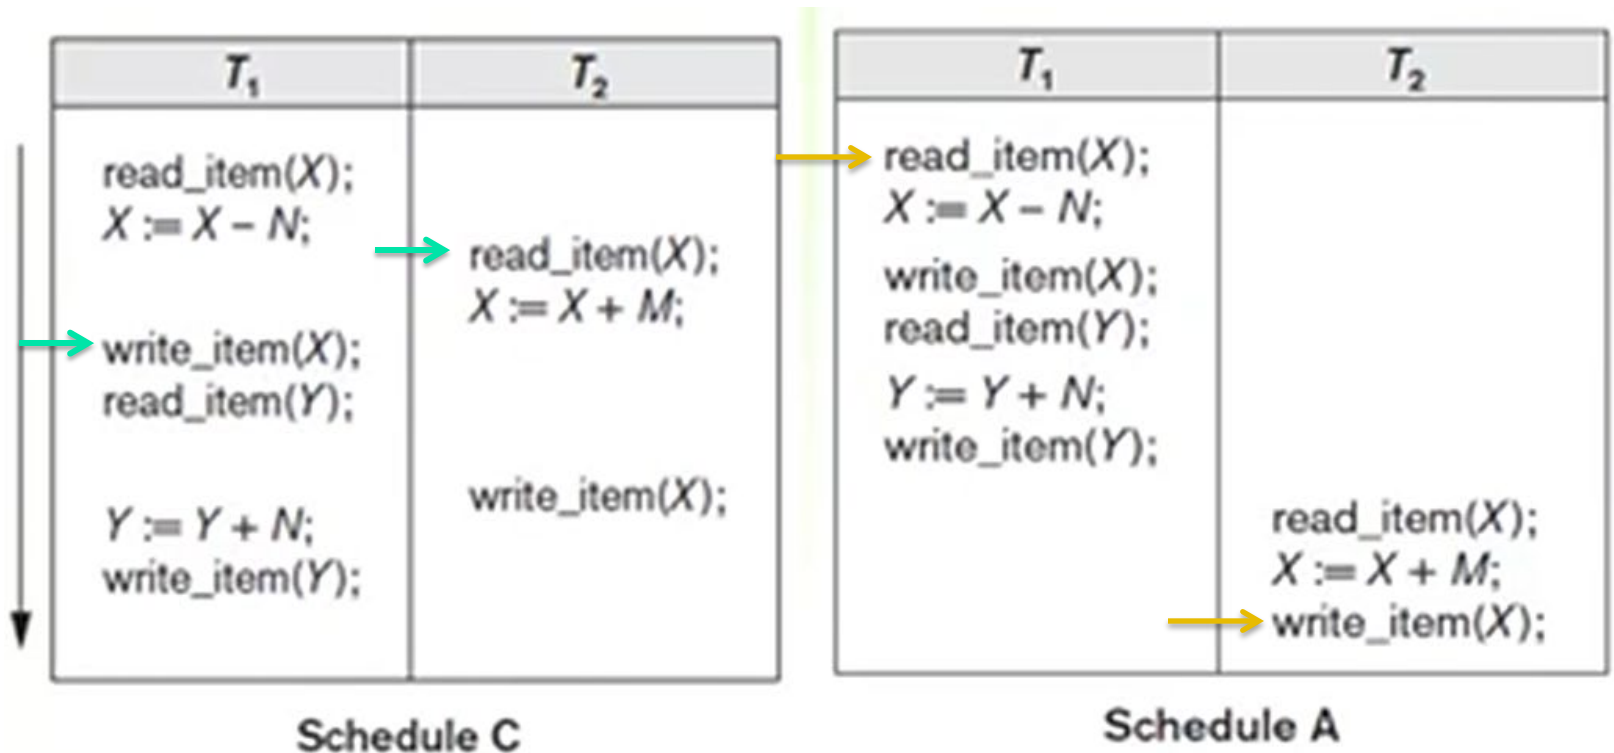
Schedule A

Schedule A=schedule D

**Conflict serializable schedule**



# Conflict serializable schedule



**Not-Conflict serializable schedule**

---

# Testing for Serializability of a Schedule

# Testing for Serializability of a Schedule

---

- There is a simple algorithm for determining whether a particular schedule is (conflict) serializable or not.
- Most concurrency control methods do *not* actually test for Serializability. Rather protocols, or **rules**, are developed that guarantee that any schedule that follows these rules will be serializable.
- Some methods **guarantee Serializability** in **most cases**, but do not guarantee it **absolutely**, in order to reduce the overhead of concurrency control.

# Testing for Serializability of a Schedule

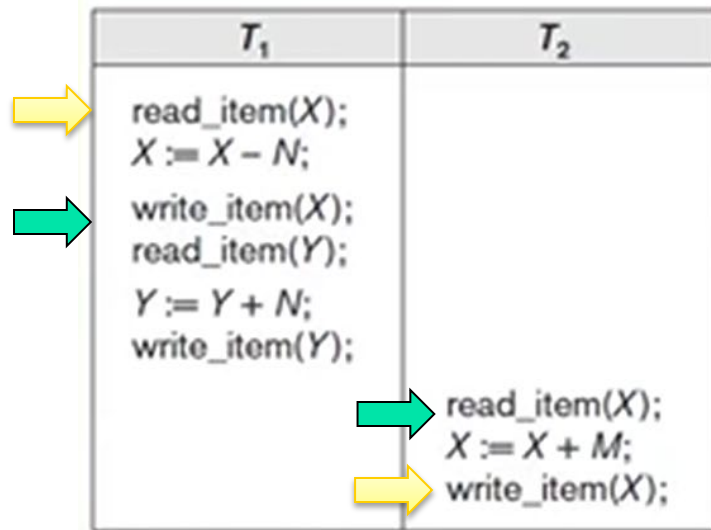
Algorithm 20.1 can be used to test a schedule for conflict serializability. The algorithm looks at only the `read_item` and `write_item` operations in a schedule to construct a **precedence graph** (or **serialization graph**), which is a **directed graph**  $G = (N, E)$  that consists of a set of nodes  $N = \{T_1, T_2, \dots, T_n\}$  and a set of directed edges  $E = \{e_1, e_2, \dots, e_m\}$ . There is one node in the graph for each transaction  $T_i$  in the schedule. Each edge  $e_i$  in the graph is of the form  $(T_j \rightarrow T_k)$ ,  $1 \leq j \leq n$ ,  $1 \leq k \leq n$ , where  $T_j$  is the **starting node** of  $e_i$  and  $T_k$  is the **ending node** of  $e_i$ . Such an edge from node  $T_j$  to node  $T_k$  is created by the algorithm if a pair of conflicting operations exist in  $T_j$  and  $T_k$  and the conflicting operation in  $T_j$  appears in the schedule *before* the *conflicting operation* in  $T_k$ .

# Testing for Serializability of a Schedule


## Algorithm 20.1. Testing Conflict Serializability of a Schedule $S$

1. For each transaction  $T_i$  participating in schedule  $S$ , create a node labeled  $T_i$  in the precedence graph.
2. For each case in  $S$  where  $T_j$  executes a `read_item(X)` after  $T_i$  executes a `write_item(X)`, create an edge  $(T_i \rightarrow T_j)$  in the precedence graph.
3. For each case in  $S$  where  $T_j$  executes a `write_item(X)` after  $T_i$  executes a `read_item(X)`, create an edge  $(T_i \rightarrow T_j)$  in the precedence graph.
4. For each case in  $S$  where  $T_j$  executes a `write_item(X)` after  $T_i$  executes a `write_item(X)`, create an edge  $(T_i \rightarrow T_j)$  in the precedence graph.
5. The schedule  $S$  is serializable if and only if the precedence graph has no cycles.

# Testing for Serializability of a Schedule

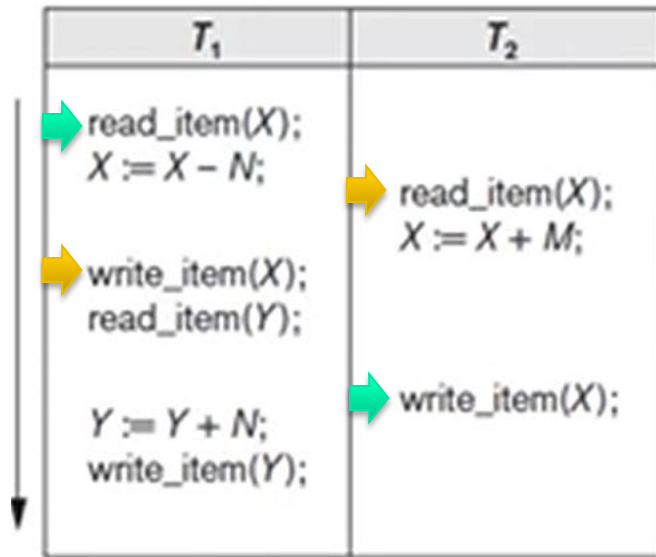


Schedule A




**No cycles,**  
Serializable schedule

# Testing for Serializability of a Schedule

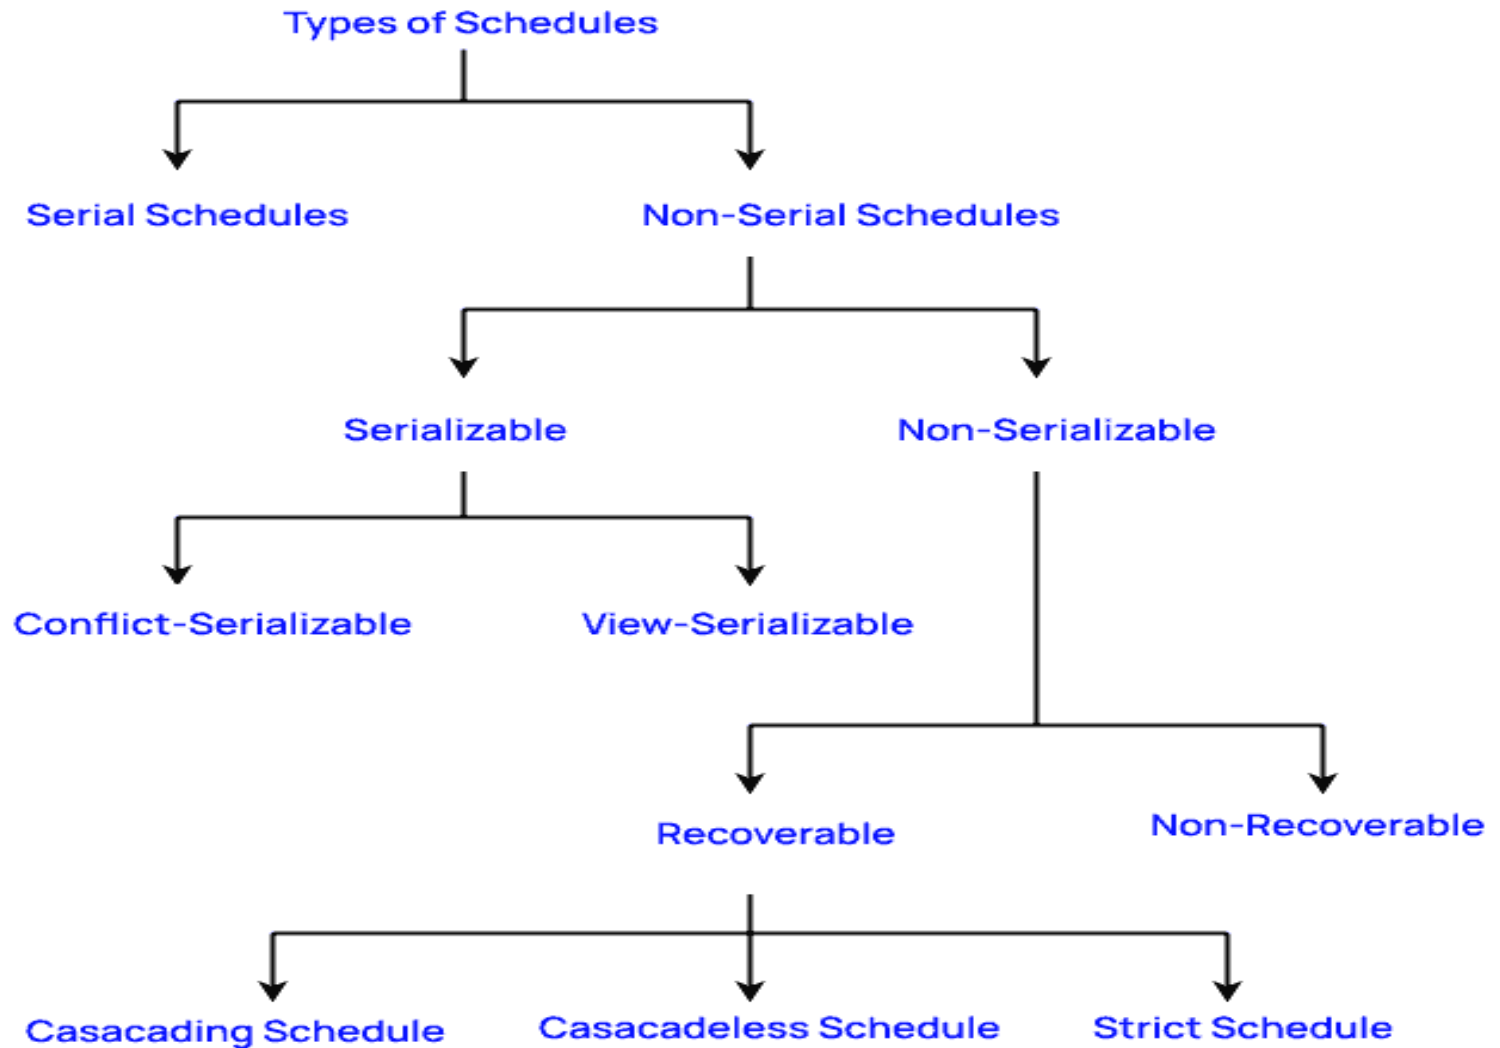


Schedule C



**Has cycles, in the same direction**  
**Non serializable**

# Types of schedules in DBMS





*Thank You!*

**Any questions?** 