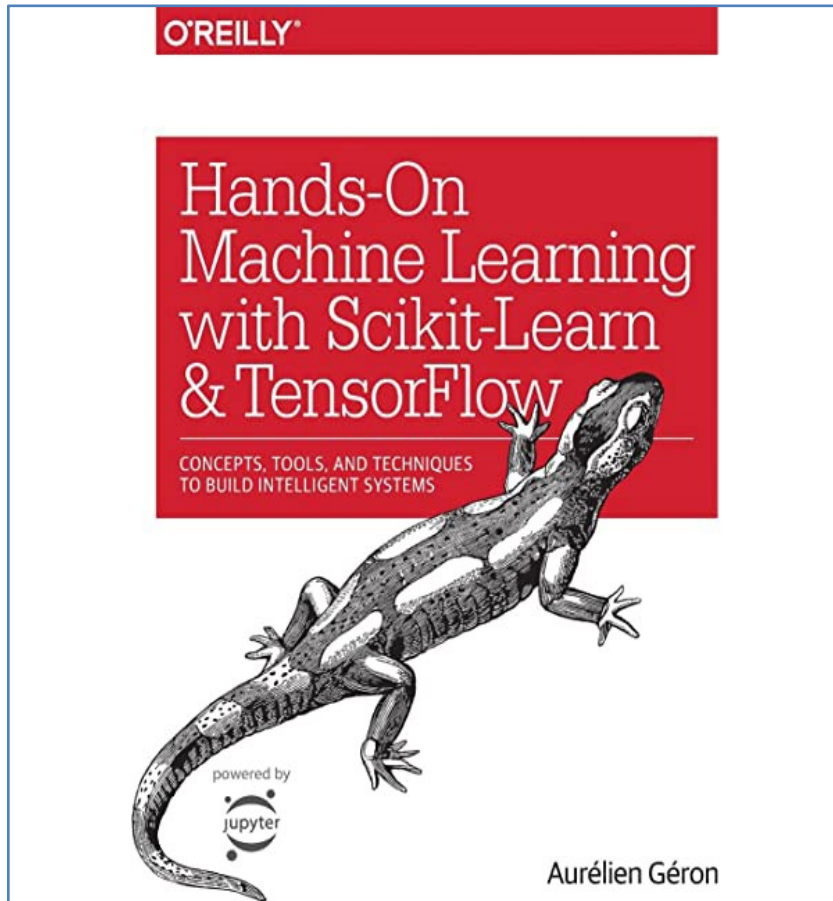


Machine Learning



Introduced by
Dr. Ebtsam Adel



Machine Learning

Chapter 3

Classification

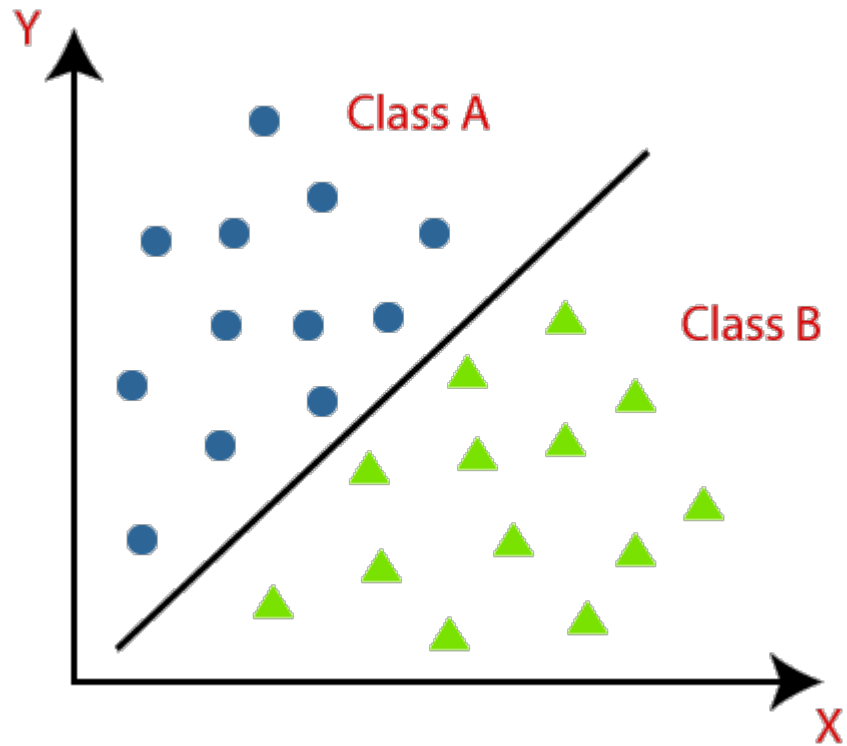


Classification

What is the Classification Algorithm?

- The Classification algorithm is a **Supervised** Learning technique that is used to identify the **category of new observations** on the basis of **training data**.
- In Classification, a program learns from the given dataset or observations and then classifies new observation into a number of **classes** or **groups**. Such as, **Yes or No, 0 or 1, Spam or Not Spam, cat or dog**, etc. Classes can be called as targets/labels or categories.

Classification



- Unlike regression, the output variable of Classification is a **category**, not a value, such as "Green or Blue", "fruit or animal", etc.

Classification

The most commonly used classifiers are as follows:

- ✓ [Logistic regression](#)
- ✓ [Naive Bayes classifier](#)
- ✓ [Support vector machines](#)
- ✓ [k-nearest neighbor](#)
- ✓ [Decision trees](#)
- ✓ [Random forests](#)
- ✓ [Neural networks](#)
- ✓

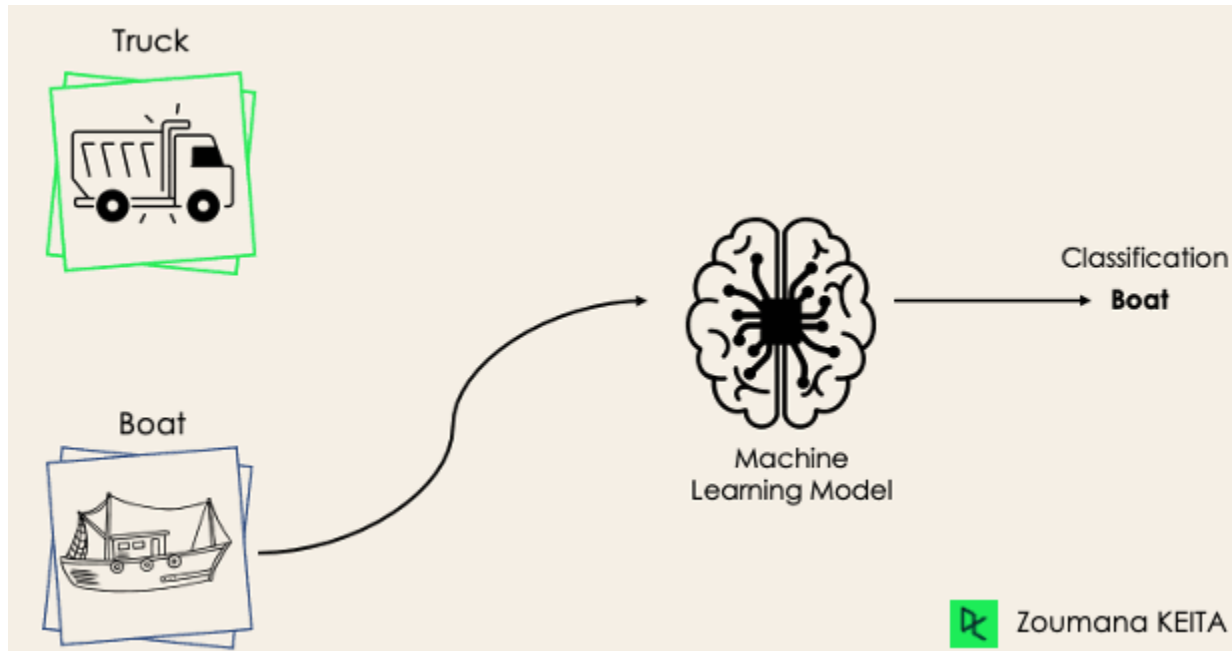
Classification

Different Types of Classification Tasks in Machine Learning

There are many main **classification tasks** in Machine learning.

- **Binary Classification**

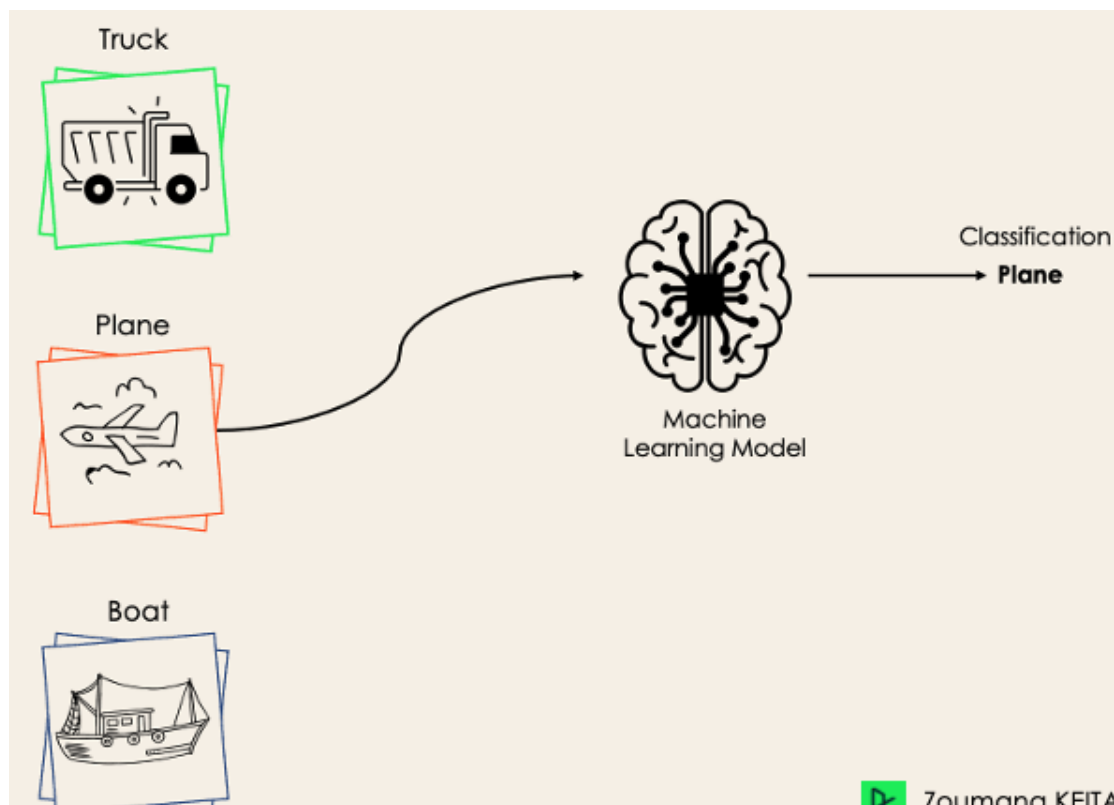
In a binary classification task, the goal is to classify the input data into **two exclusive categories**.



Classification

- **Multi-Class Classification**

The multi-class classification, on the other hand, has **at least two exclusive** class labels, where the goal is to predict to which class a given input example belongs to.



Classification

MNIST

- the MNIST dataset, which is a set of 70,000 small images of digits handwritten by high school students and employees of the US Census Bureau.
- Each image is labeled with the digit it represents.
- This set has been studied so much that it is often called the “Hello World” of Machine Learning.
- There are 70,000 images, and each image has 784 features. This is because each image is 28×28 pixels, and each feature simply represents one pixel’s intensity, from 0 (white) to 255 (black).

Classification

- The MNIST dataset is actually already split into a **training** set (the first **60,000 images**) and a **test** set (the last **10,000 images**):

```
X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]
```



Classification

```
>>> from sklearn.datasets import fetch_openml
>>> mnist = fetch_openml('mnist_784', version=1)
>>> mnist.keys()
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'details',
           'categories', 'url'])
```

- A **DESCR** key describing the dataset.
- A **data** key containing an array with one row per instance and one column per feature.
- A **target** key containing an array with the labels.

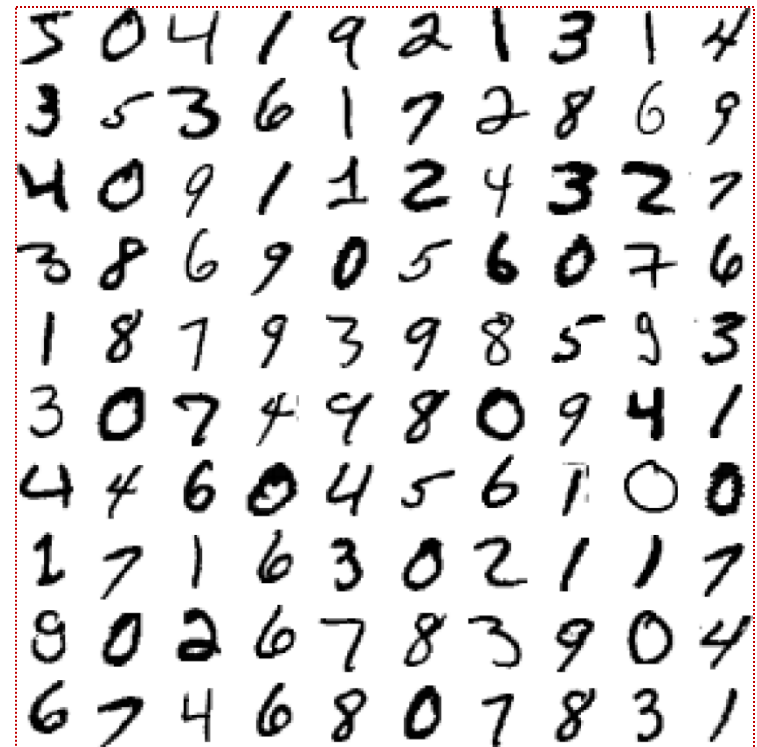
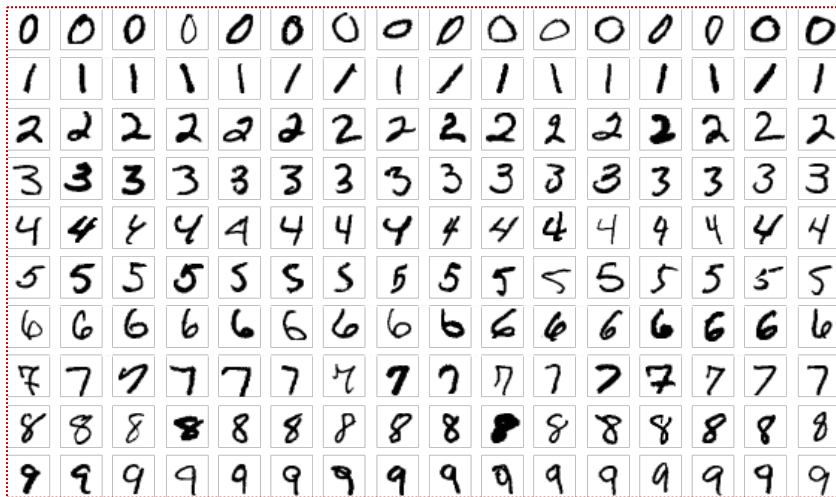
Classification

the **complexity** of the classification task.



Classification

- The training set is already **shuffled** for us, which is good as this guarantees that all **cross-validation folds** will be **similar**. Moreover, some learning algorithms **are sensitive to the order of the training instances**, and they perform poorly if they get **many similar instances** in a row.
- **Shuffling** the dataset ensures that this won't happen.

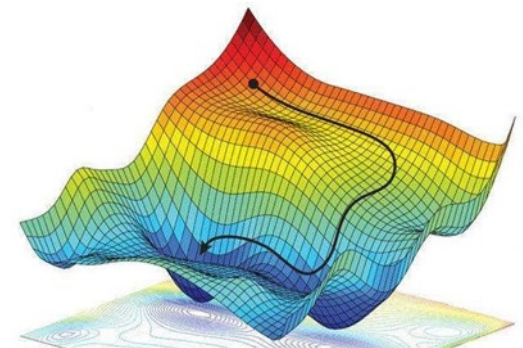


Training a Binary Classifier

- This “**5-detector**” will be an example of a **binary classifier**, capable of separating between just **two classes, 5 and not 5**.
- Let's create the target vectors for this classification task:

```
y_train_5 = (y_train == 5) #      True for all 5s, False for all other digits.  
y_test_5 = (y_test == 5)
```

- good place to start is with a **Stochastic Gradient Descent (SGD)** classifier, using Scikit-Learn's **SGDClassifier** class.
- This classifier has the advantage of being capable of handling **very large datasets** efficiently. This is in part because SGD deals with training instances independently, **one at a time**.
- Suitable for **online learning**.



Performance Measures

Machine Learning - Performance Metrics

Performance metrics in machine learning are used to **evaluate** the **performance** of a machine learning model.

- **Performance metrics** are important because they help us understand how well our model is **performing** and whether it is **meeting our requirements**.
- There are many performance metrics that can be used in machine learning, depending on the **type of problem being solved** and the **specific requirements** of the problem.
- Some common performance metrics include: **Accuracy** , **Recall**, **F1 Score**,...

Machine Learning - Performance Metrics

Confusion Matrix – A confusion matrix is a **table** that is used to evaluate the performance of a classification model. It shows the number of **true positives**, **true negatives**, **false positives**, and **false negatives** for each class in the dataset.

Confusion Matrix

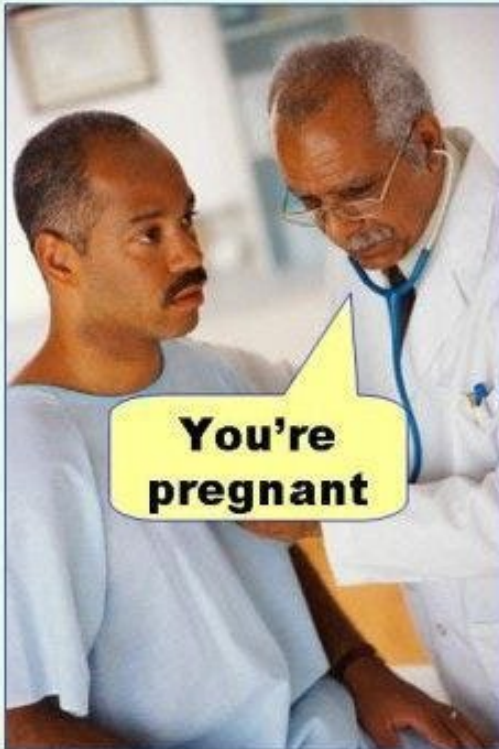
	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	True Positives (TPs)	False Positives (FPs)
Predicted Negative (0)	False Negatives (FNs)	True Negatives (TNs)

FN: Ham mail in spam folder.

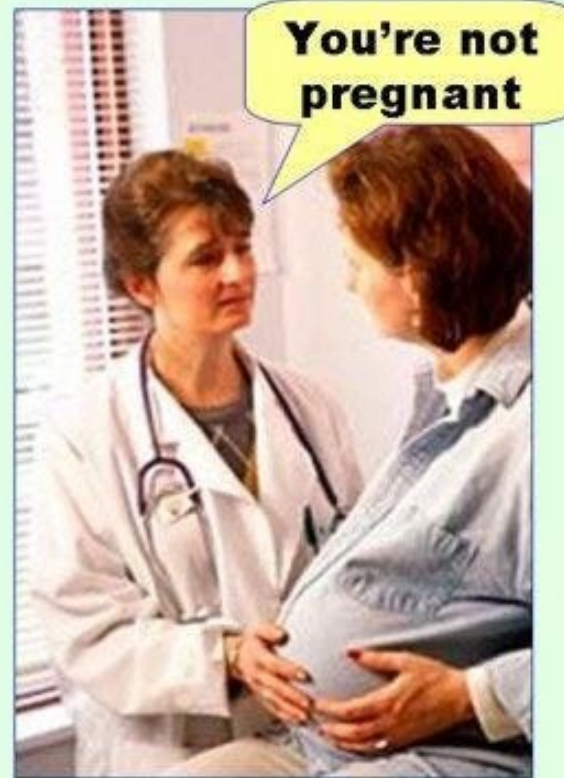
FP: spam email in inbox.

Machine Learning - Performance Metrics

Type I error
(false positive)



Type II error
(false negative)



Machine Learning - Performance Metrics

- True Positive(**TP**): the prediction outcome is true, and it is true in reality, also.
- True Negative(**TN**): the prediction outcome is false, and it is false in reality, also.
- False Positive(**FP**): prediction outcomes are true, but they are false in actuality.
- False Negative(**FN**): predictions are false, and they are true in actuality.

	Actual YES	Actual NO
Predicted YES	True Positive	False Positive
Predicted NO	False Negative	True Negative

Machine Learning - Performance Metrics

- **Accuracy** – Accuracy is one of the most basic performance metrics and measures the proportion of correctly classified instances in the dataset.
- It is calculated as the number of **correctly classified instances** divided by the **total number of instances** in the dataset.
- **Precision** :It is calculated as the number of **true positive** instances divided by the **sum of true positive and false positive** instances.

$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Total number of predictions}}$$

$$\text{Accuracy} = \frac{TP+TN}{TP+TN+FP+FN}$$

$$\text{Precision} = \frac{TP}{(TP + FP)}$$

Machine Learning - Performance Metrics

- **Recall** “sensitivity”: It is calculated as the number of true positive instances divided by the sum of true positive and false negative instances.

$$\text{Recall} = \frac{TP}{TP + FN}$$

Machine Learning - Performance Metrics

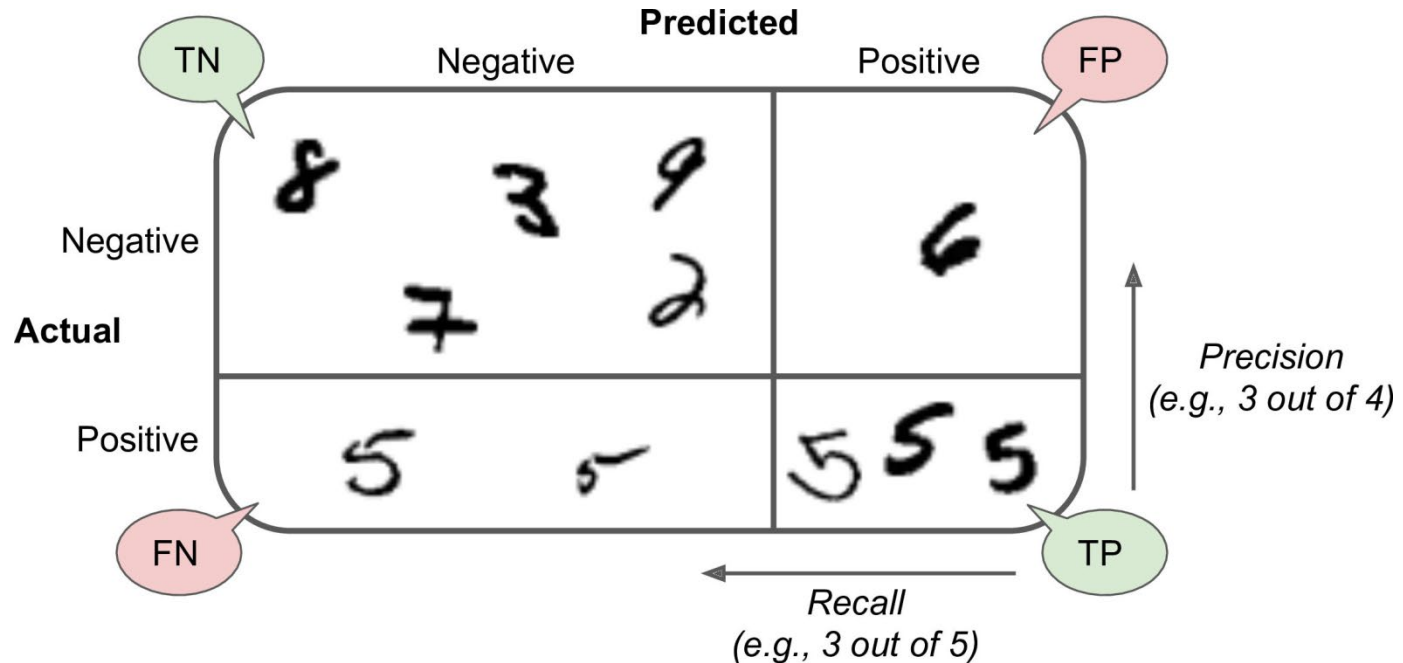
- **F1 Score** – F1 score is the harmonic mean of precision and recall. It is a balanced measure that takes into account both precision and recall.
- It is calculated as $2 * (\text{precision} \times \text{recall}) / (\text{precision} + \text{recall})$.
- the classifier will only get a high F1 score if both recall and precision are high.

$$F_1 = \frac{2}{\frac{1}{\text{precision}} + \frac{1}{\text{recall}}} = 2 \times \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}} = \frac{TP}{TP + \frac{FN + FP}{2}}$$

To compute the F_1 score, simply call the `f1_score()` function:

```
>>> from sklearn.metrics import f1_score
>>> f1_score(y_train_5, y_train_pred)
0.7420962043663375
```

Machine Learning - Performance Metrics



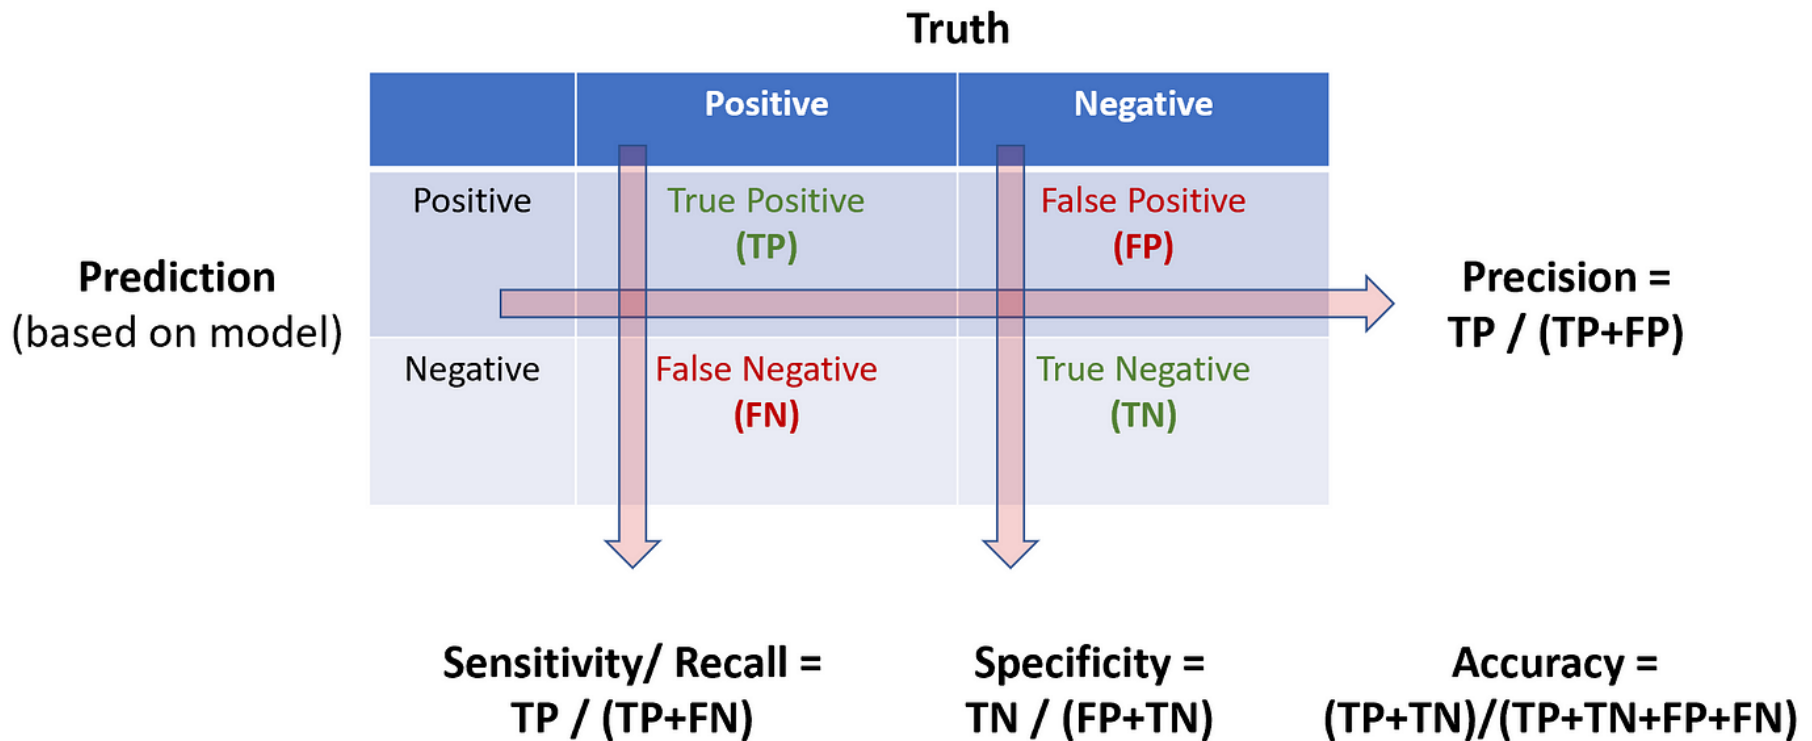
Equation 3-1. Precision

$$\text{precision} = \frac{TP}{TP + FP}$$

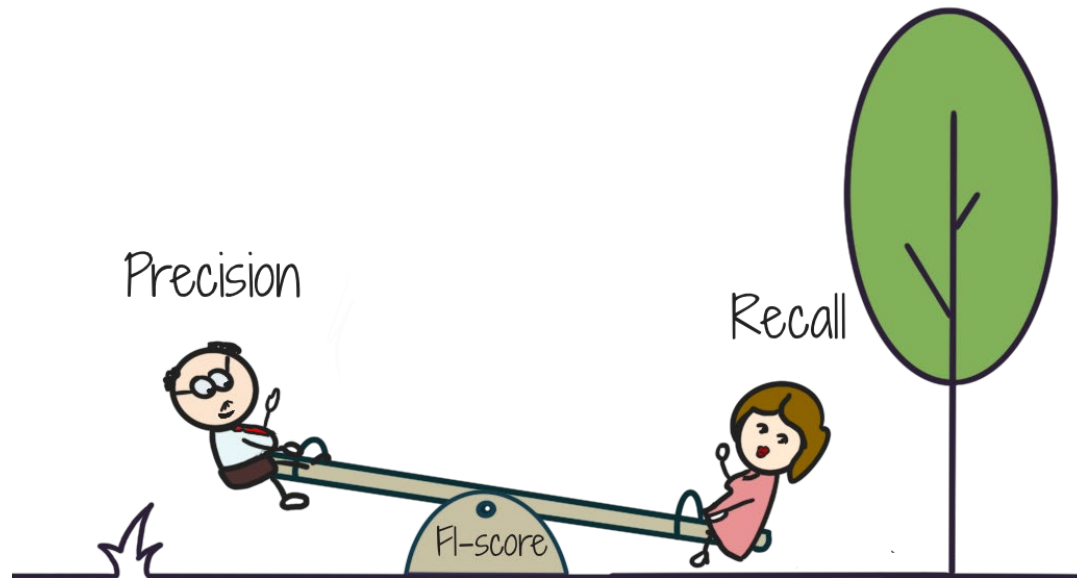
Equation 3-2. Recall

$$\text{recall} = \frac{TP}{TP + FN}$$

Machine Learning - Performance Metrics



Precision/Recall Tradeoff

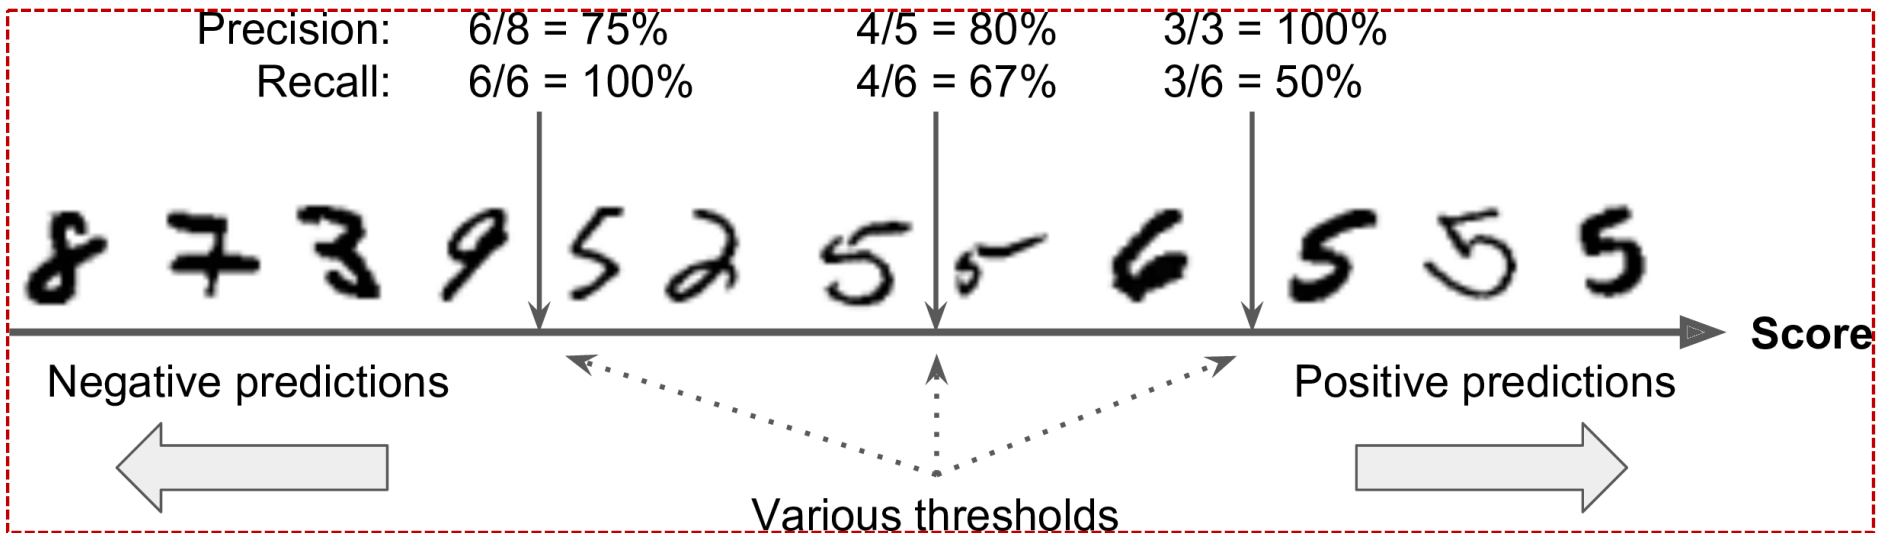


Machine Learning - Precision/Recall Tradeoff

- Unfortunately, you can't have it both ways: **increasing precision reduces recall**, and vice versa. This is called the **precision/recall tradeoff**.
- To understand this **tradeoff**, let's look at how the **SGDClassifier** makes its classification decisions. For each instance, it computes a score based on **a decision function**, and if that score is greater than a **threshold**, it assigns the instance to the positive class, or else it assigns it to the negative class.



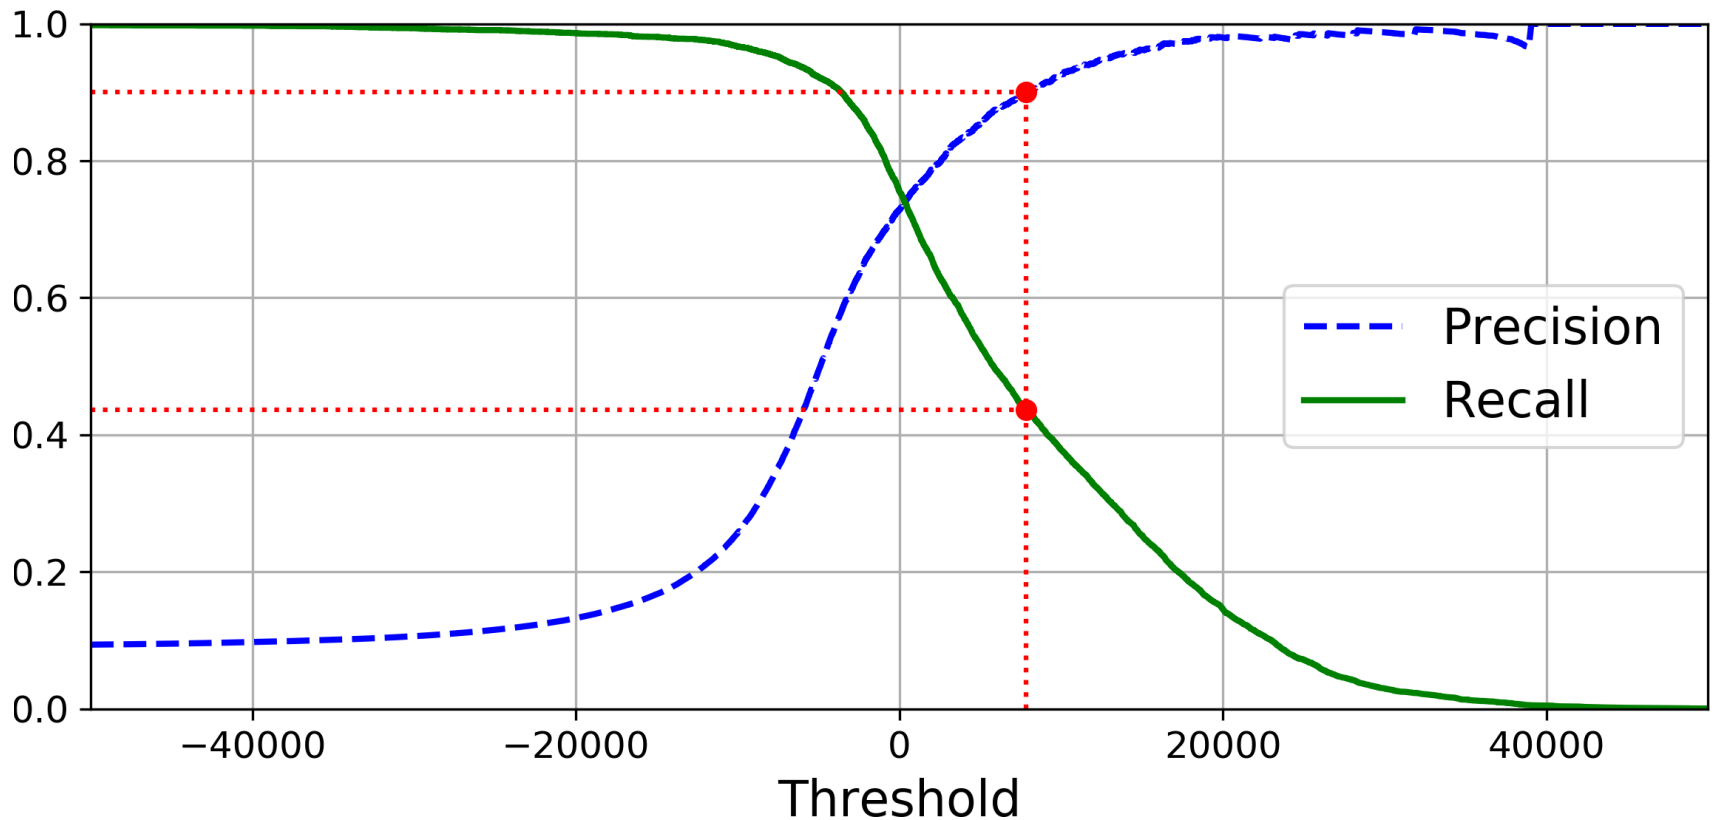
Machine Learning - Precision/Recall Tradeoff



✓ Raising threshold increases precision, but decrease recall.



Machine Learning - Precision/Recall Tradeoff



- If someone says “let’s reach 99% precision,” you should ask, “at what recall?”
- *Determine threshold from the curve.*

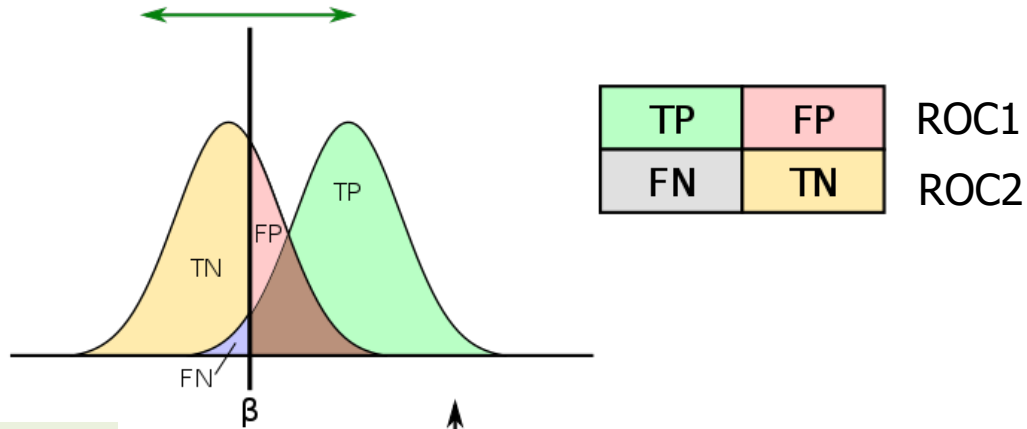
The ROC Curve

Receiver Operating Characteristic

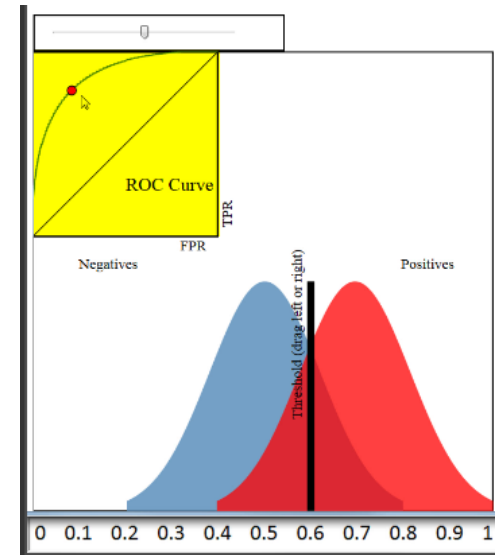
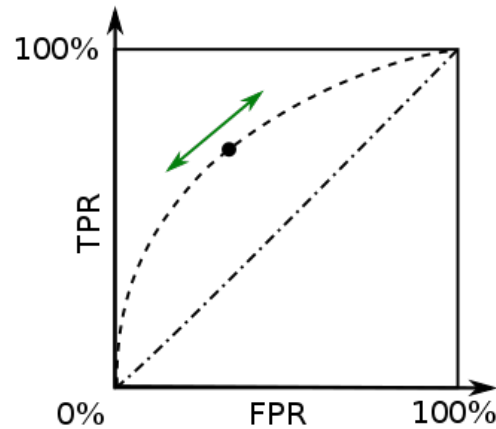
Machine Learning - ROC Score

- The **receiver operating characteristic** (**ROC**) curve is another common tool used with **binary** classifiers.
- It is very similar to the **precision/recall curve**, but instead of plotting precision versus recall, the ROC curve plots the **true positive** rate (another name for recall) against the **false positive** rate.
- The **FPR** is the ratio of negative instances that are incorrectly classified as positive. It is **equal to one minus** the **true negative rate**, which is the ratio of **negative instances** that are correctly classified as **negative**.
- The **TNR** is also called **specificity**. Hence the ROC curve plots sensitivity (recall) versus $1 - \text{specificity}$

Machine Learning - ROC Score



histogram

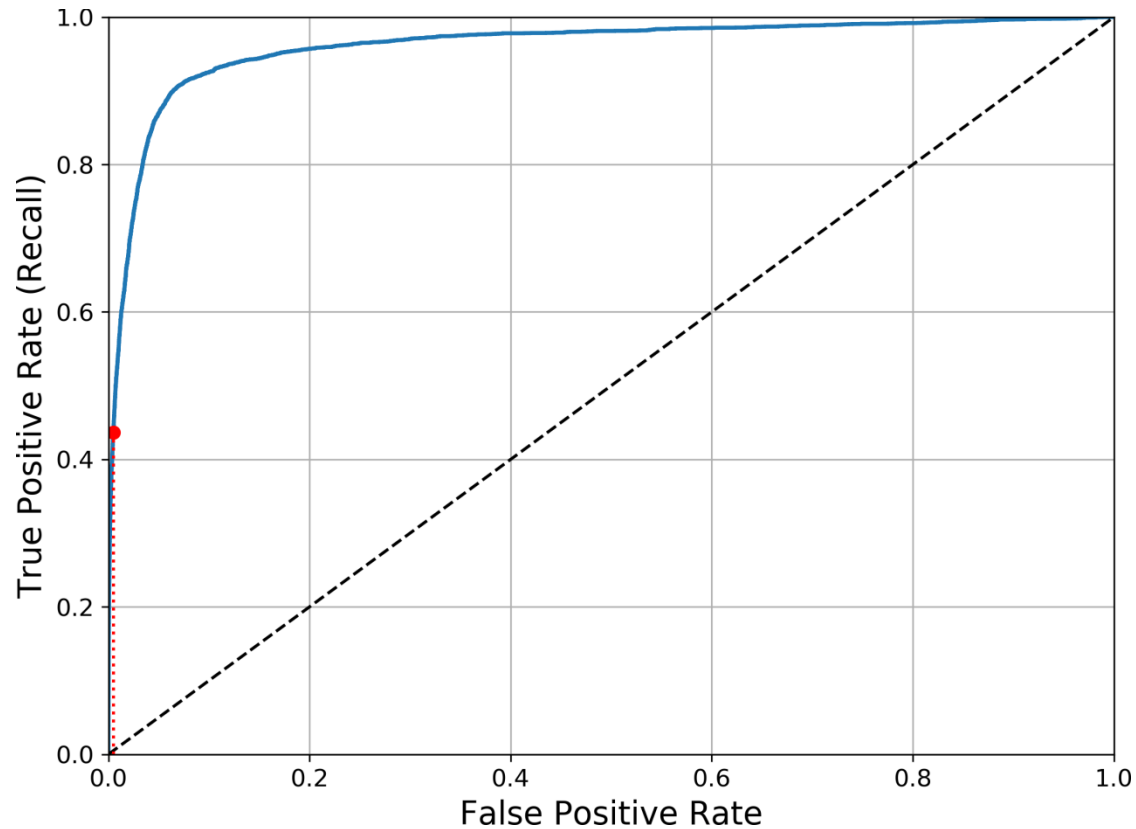


Machine Learning - ROC Score

What is the ROC Curve?

- The ROC curve is a graphical representation of the performance of a binary classification model at different threshold values.
- It plots the true positive rate (**TPR**) on the y-axis and the false positive rate (**FPR**) on the x-axis.
- The **TPR** is the proportion of actual positive cases that are correctly identified by the model, while the **FPR** is the proportion of actual negative cases that are incorrectly classified as positive by the model.

Machine Learning - ROC Score



Machine Learning - ROC AUC Score

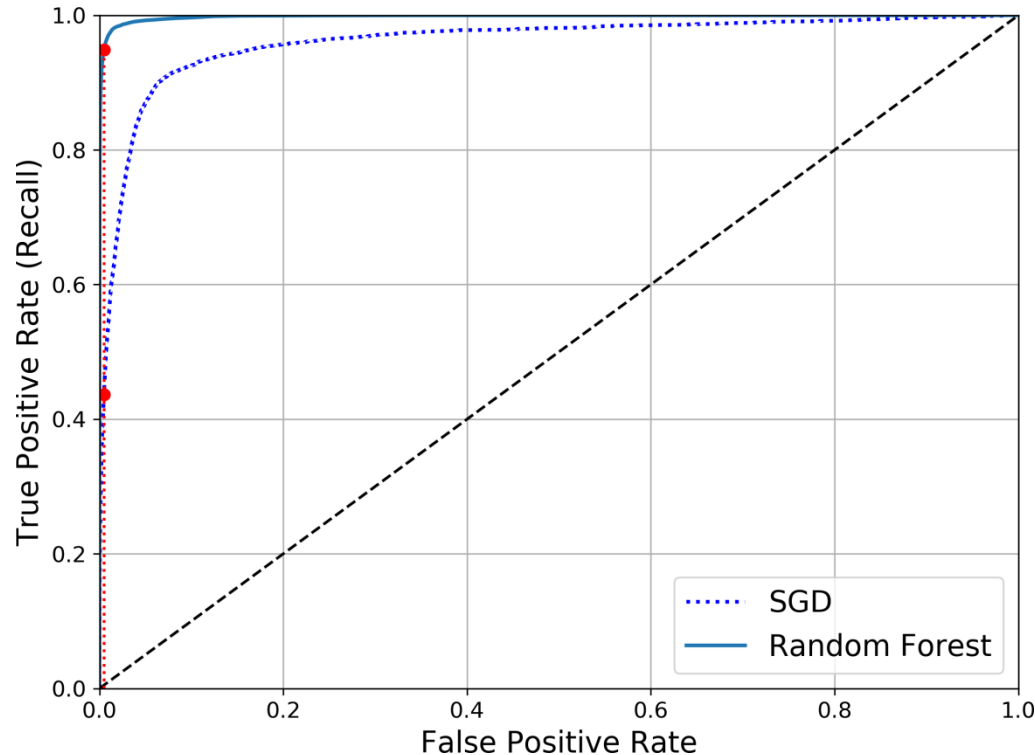
- One way to compare classifiers is to measure the **area under the curve (AUC)**. A perfect classifier will have a **ROC AUC equal to 1**, whereas a purely random classifier will have a **ROC AUC equal to 0.5**.
- Scikit-Learn provides a function to compute the ROC AUC:

```
>>> from sklearn.metrics import roc_auc_score  
>>> roc_auc_score(y_train_5, y_scores)  
0.9611778893101814
```

Why is the AUC-ROC Curve Important?

The AUC-ROC curve is an important performance metric in machine learning because it provides a comprehensive measure of a model's ability to recognize **positive and negative cases**.

Machine Learning - ROC AUC Score



- Let's train a **RandomForestClassifier** and compare its ROC curve and ROC AUC score to the **SGDClassifier**.
- the **RandomForestClassifier's** ROC curve looks much **better** than the **SGDClassifier's**.

Practical Part

Practical part

- Calculate all the performance measurements in Python
- Implement the AUC ROC Curve in Python

Thank You!

Any questions? 